

Rethinking the policy optimization with Gradient Estimators

Dongkyu Derek Cho, Thomas Voice

Abstract

Permissive policy optimization can learn rapidly from rollout data and then fail abruptly after many successful updates. Policy-deviation bounds explain why a large change can be unsafe, but they do not say when a finite rollout stops supporting an accurate update for the current policy. We study this transition through two ordered questions. First, when is the ordinary full-trajectory importance-sampling (IS) estimator still reliable? We show that normalized effective sample size (ESS) controls the policy-mismatch component of its finite-sample error and, through smoothness, yields a high-probability one-step improvement guarantee. Second, when does coefficient control provide a better update? Lower gradient-estimation MSE is not enough: the controlled estimator must also retain sufficient alignment with the population gradient. We capture this requirement with a full fixed-step improvement certificate. An exactly solvable Gaussian bandit shows that the two criteria can switch at different ESS levels: PPO can have lower MSE while IS still gives the larger exact expected improvement. Along the example’s monotone off-policy path, the safe certificate oracle—IS, PPO, or no update—reduces exactly to two ESS thresholds. In the reported large-model trajectories, ESS-conditioned safeguards preserve early learning while preventing late failure.

1 Introduction

Reinforcement learning with verifiable rewards repeatedly collects responses from a rollout policy and reuses each batch for one or more policy updates. A permissive update can work surprisingly well in this setting: performance rises quickly and remains stable, but then collapses abruptly. Figure 1 shows this delayed failure. Rollout data support useful learning for many updates before the same update rule becomes unstable.

Classical theory explains why moving too far from the rollout policy can be dangerous [11, 13, 14]. That explanation is essential, but it does not explain why failure is delayed. A policy-deviation bound controls a population approximation; it does not track how much information a finite rollout still provides for the gradient of the current policy. The practical question raised by Figure 1 is therefore statistical: when does the rollout cease to support a reliable next update?

Our starting point is that nominal batch size can remain fixed while effective support contracts. Importance weights become increasingly uneven as the current policy moves away from the rollout policy, so fewer responses determine the update. Normalized ESS measures this contraction. We connect it first to gradient reliability and then to policy improvement.

This separation leads to two ordered questions. First, when is the IS importance-weighted gradient reliable enough to improve the population objective? Second, if IS becomes unreliable, when is a controlled estimator actually better? The second question cannot be answered by ESS or MSE alone. Clipping may reduce variance and MSE while also removing useful population-gradient alignment. We therefore compare candidate estimators using the full fixed-step improvement certificate.




Figure 1: A delayed failure in Qwen3-30B-A3B training. The permissive runs improve for many updates and fail only after the effective support of the rollout data has sharply deteriorated. This transition is the phenomenon studied in the paper.

An exactly solvable Gaussian bandit makes the distinction concrete. Every population moment and the expected one-step gain are available without Monte Carlo error, allowing us to separate the MSE crossover from the actual improvement crossover. We then test the resulting regime prediction in large-model training. These experiments use ESS as a diagnostic of effective support and test the predicted support-to-intervention ordering.

1.1 Contributions

1. We show how normalized ESS controls the finite-sample error of the IS gradient estimator from a fixed rollout and yields a one-step population improvement guarantee.
2. We distinguish estimator accuracy from update quality and derive the fixed-step condition under which coefficient control has the stronger improvement certificate.
3. We separate the two crossovers exactly in a Gaussian bandit and then evaluate ESS-conditioned safeguards in large-model training.

2 Related work

Policy-deviation guarantees. Trust-region and proximal methods control policy movement to obtain stable improvement guarantees or practical update rules [11, 13, 14]. These analyses provide the standard explanation for why large policy changes can be unsafe. Our framework is complementary: it studies when a fixed finite rollout ceases to provide a reliable estimate of the current policy gradient. The distinction is between a population-level approximation guarantee and a finite-sample reliability question.

Importance sampling and effective support. Importance-sampling policy optimization uses likelihood-ratio moments and Rényi divergence to quantify the reliability of off-policy estimators [5, 6]. In these results, effective rather than nominal sample size controls estimation error. We apply the same principle directly to the vector-valued policy-gradient estimator and use normalized ESS as the regime coordinate.

Reliable policy-gradient updates. Smoothness-based analyses translate the error of a finite gradient estimate into a lower bound on policy improvement and use this relation to choose step sizes or sample sizes [8, 10]. We combine this optimization argument with an ESS-dependent off-policy reliability bound.

Stability in RLVR. Recent methods alter ratio boundaries, sequence weights, divergence controls, or update rules to improve the stability of large-model reinforcement learning [2, 3, 7, 9, 16, 18, 19]. Rather than proposing another static boundary, we study the statistical transition that makes a permissive update unreliable and then ask when an intervention can recover reliability.

3 Preliminaries

This section provides the background and notation used in the following analysis.

3.1 Policy-gradient estimators from a fixed rollout

Problem formulation. Let $X \sim \nu$ denote a context, $Q(\cdot | X)$ the rollout policy, and $P_\theta(\cdot | X)$ the current policy. A complete sample is $Z = (X, Y)$, where $Y = (Y_1, \dots, Y_T)$ and $R(Z)$ is its reward. For each factor t , define the likelihood ratio and score

$$r_{\theta,t}(Z) := \frac{P_\theta(Y_t | X, Y_{<t})}{Q(Y_t | X, Y_{<t})}, \quad \ell_{\theta,t}(Z) := \nabla_\theta \log P_\theta(Y_t | X, Y_{<t}). \quad (1)$$

Here t indexes a policy factor, such as a token or decision step. Because both policies factorize along the trajectory, the complete-sample likelihood ratio is

$$W_\theta(Z) := \frac{P_\theta(Y | X)}{Q(Y | X)} = \prod_{t=1}^T r_{\theta,t}(Z). \quad (2)$$

The learning objective is to maximize the expected reward $J(\theta) := \mathbb{E}_{P_\theta}[R]$. Let $A(Z) := R(Z) - b(X)$ be a baseline-centered reward, where $b(X)$ does not depend on Y . Under $P_\theta(\cdot | X) \ll Q(\cdot | X)$, change of measure expresses the population gradient under the rollout distribution:

$$g(\theta) := \nabla J(\theta) = \mathbb{E}_Q[W_\theta f_\theta], \quad f_\theta(Z) := A(Z) \sum_{t=1}^T \ell_{\theta,t}(Z) = A(Z) \nabla_\theta \log P_\theta(Y | X). \quad (3)$$

While trust-region arguments motivate safeguards such as TRPO and PPO-style clipping [13, 14], these arguments are typically made at the population level. Practical algorithms, however, do not observe $g(\theta)$ and must rely on its finite-sample counterpart. Given i.i.d. rollout samples $Z_i \sim \nu(dx)Q(dy | x)$, $i = 1, \dots, N$, the ordinary full-trajectory importance-sampling (IS) policy-gradient estimator is

$$\hat{g}_N(\theta) := \frac{1}{N} \sum_{i=1}^N W_{\theta,i} f_{\theta,i}. \quad (4)$$

The importance-weighted nature of this estimator can produce poor agreement between the finite-sample update and the population gradient [5, 6]. This exposes an important but often overlooked question: when is the finite-sample gradient reliable?

Choice of gradient estimator. To formalize this comparison, write $\mathbf{r}_\theta = (r_{\theta,1}, \dots, r_{\theta,T})$ and let $u = (u_1, \dots, u_T)$ be a collection of coefficient rules, with $u_t : \mathbb{R}_+^T \times \mathbb{R} \rightarrow \mathbb{R}$. Define

$$\psi_\theta(u; Z) := A(Z) \sum_{t=1}^T u_t(\mathbf{r}_\theta, A(Z)) \ell_{\theta,t}(Z), \quad \hat{g}_N(u; \theta) := \frac{1}{N} \sum_{i=1}^N \psi_\theta(u; Z_i). \quad (5)$$

Existing ratio-based policy-gradient estimators can be represented by particular choices of u . The IS estimator sets $u_{\text{IS},t} = W_\theta$ for every factor. One popular alternative is PPO/GRPO ratio clipping, whose clipped surrogate induces the following gradient-masking rule with radius $0 < \epsilon < 1$ [14, 15]:

$$u_{\text{P},t}(\mathbf{r}, A) := r_t [\mathbf{1}\{A \geq 0, r_t \leq 1 + \epsilon\} + \mathbf{1}\{A < 0, r_t \geq 1 - \epsilon\}]. \quad (6)$$

Such clipping rules create a conventional bias–variance tradeoff. More aggressive clipping can reduce sampling variation, but it can also increase bias by removing gradient terms. This raises the natural question: what governs the reliability of a gradient estimator, and when is clipping worth the signal it removes? We answer the first part by introducing Rényi divergence and effective sample size (ESS).

3.2 Gradient estimators and Rényi divergence

The dispersion of the importance weights indicates how dissimilar P_θ and Q are. Rényi divergence formalizes this notion as an information-theoretic dissimilarity between probability measures [12, 17], while importance-sampling analyses connect it to the finite-sample behavior of importance-weighted means [5, 6].

Rényi divergence is a family of measures indexed by a moment order $\alpha > 1$. For the current and rollout policies, define

$$D_\alpha(P_\theta \| Q) := \frac{1}{\alpha - 1} \log \mathbb{E}_Q[W_\theta^\alpha], \quad d_\alpha(P_\theta \| Q) := \exp\{D_\alpha(P_\theta \| Q)\}. \quad (7)$$

Larger values of α place progressively more emphasis on rare, large likelihood ratios. The order-two case is central because $\mathbb{E}_Q[W_\theta^2]$ controls standard MSE bounds for bounded importance-weighted contributions. Importantly, its reciprocal is the normalized effective sample size [5, 6]:

$$\rho_\theta := \frac{1}{d_2(P_\theta \| Q)} = \frac{1}{\mathbb{E}_Q[W_\theta^2]} = \exp\{-D_2(P_\theta \| Q)\}. \quad (8)$$

When $\mathbb{E}_Q[W_\theta^2] < \infty$, Jensen’s inequality and $\mathbb{E}_Q[W_\theta] = 1$ give $\rho_\theta \in (0, 1]$, and $N\rho_\theta$ is the effective sample count corresponding to nominal batch size N . As the second moment of the likelihood ratio grows, both ρ_θ and this effective count contract. The identity $\mathbb{E}_Q[(W_\theta - 1)^2] = \rho_\theta^{-1} - 1$ shows that ρ_θ is exactly the inverse second-moment measure of ratio dispersion.

Importance-sampling theory makes this quantity operational. If $\|f_\theta\|_\infty := \sup_z \|f_\theta(z)\|_2 < \infty$, the MSE of the IS estimator is at most $\|f_\theta\|_\infty^2 / (N\rho_\theta)$. Normalized ESS therefore captures the policy-mismatch component of IS reliability. The next section determines when that reliability is sufficient for improvement and then asks whether coefficient control offers a better update.

4 Theoretical Analysis

This section answers the preceding questions in order. We first connect gradient-estimator reliability to policy improvement, then use normalized ESS to control the reliability of the IS estimator, and finally ask when clipping provides the stronger update.

4.1 Gradient-estimator reliability and policy improvement

The preceding section distinguishes the finite-sample IS estimate from the population gradient it targets. Here, we connect estimator reliability to a lower bound on policy improvement. The following lemma formalizes this connection for an arbitrary gradient estimator.

Lemma 1 (A reliable gradient estimate yields policy improvement). *Let $\hat{g}$ be an estimate of $g = \nabla J(\theta)$ satisfying $\|\hat{g} - g\|_2 \leq \varepsilon$. Suppose J is L -smooth along the update segment for $L > 0$, and define $\gamma = L^{-1}(1 - \varepsilon/\|\hat{g}\|_2)_+$, with $\gamma = 0$ when $\hat{g} = 0$. Then*

$$J(\theta + \gamma\hat{g}) - J(\theta) \geq \frac{1}{2L} (\|\hat{g}\|_2 - \varepsilon)_+^2. \quad (9)$$

This deterministic lemma uses the same smoothness mechanism as existing policy-gradient analyses [8, 10]. Its implication is direct: the error radius measures how far the finite-sample direction can lie from the population gradient. When this radius is smaller than $\|\hat{g}\|_2$, the estimator yields a positive lower bound on policy improvement.

ESS and the IS update. However, the population gradient is unknown, so this reliability cannot be evaluated directly. Normalized ESS provides the bridge: for the IS estimator, it converts likelihood-ratio dispersion into a bound on gradient-estimation error. At update k , condition on the preceding updates and write $W_k := W_{\theta_k}$, $f_k := f_{\theta_k}$, $g_k := g(\theta_k)$, $\hat{g}_k := \hat{g}_N(\theta_k)$, and $\rho_k := \rho_{\theta_k}$. Assume that the N rollout samples are independent draws from Q and that $\|f_k\|_\infty := \sup_z \|f_k(z)\|_2 < \infty$.

Because the IS estimator is unbiased, its exact mean-squared error is

$$\mathbb{E}\|\hat{g}_k - g_k\|_2^2 = \frac{1}{N} \{ \mathbb{E}_Q[W_k^2 \|f_k\|_2^2] - \|g_k\|_2^2 \}. \quad (10)$$

Using $\|f_k(z)\|_2 \leq \|f_k\|_\infty$ and $\rho_k^{-1} = \mathbb{E}_Q[W_k^2]$ gives the interpretable upper bound

$$\mathbb{E}\|\hat{g}_k - g_k\|_2^2 \leq \frac{\|f_k\|_\infty^2}{N\rho_k}. \quad (11)$$

Likelihood-ratio dispersion therefore enters through the effective count $N\rho_k$, replacing the nominal sample size N . Equation (11) controls the mean-squared error; the next lemma converts this result into a high-probability error radius for the realized update.

Lemma 2 (Normalized ESS controls the gradient-error radius). *Under the preceding setup, suppose $d_2(P_{\theta_k} \| Q) < \infty$ and $\|f_k\|_\infty < \infty$. Then, for any $0 < \delta \leq 1$, with probability at least $1 - \delta$,*

$$\|\hat{g}_k - g_k\|_2 \leq \frac{\|f_k\|_\infty}{\sqrt{\delta N \rho_k}}. \quad (12)$$

The radius grows as $(N\rho_k)^{-1/2}$. Thus, low ESS weakens the reliability guarantee for the IS gradient, consistent with the motivating example in the introduction. We now combine this radius with Lemma 1 to obtain an ESS-dependent improvement guarantee.

Theorem 3 (ESS-dependent improvement for the IS estimator). *Under the conditions of Lemma 2, suppose that J is L_k -smooth along the update segment for $L_k > 0$. For any $0 < \delta \leq 1$, define*

$$\varepsilon_k(\delta) := \frac{\|f_k\|_\infty}{\sqrt{\delta N \rho_k}}, \quad \gamma_k := \frac{1}{L_k} \left(1 - \frac{\varepsilon_k(\delta)}{\|\hat{g}_k\|_2} \right)_+, \quad (13)$$

with $\gamma_k = 0$ when $\hat{g}_k = 0$. Then, with probability at least $1 - \delta$,

$$J(\theta_k + \gamma_k \hat{g}_k) - J(\theta_k) \geq \frac{1}{2L_k} (\|\hat{g}_k\|_2 - \varepsilon_k(\delta))_+^2. \quad (14)$$

Theorem 3 makes the role of ESS explicit. Its lower bound is positive when

$$\|\hat{g}_k\|_2 > \frac{\|f_k\|_\infty}{\sqrt{\delta N \rho_k}}.$$

Thus, the IS update remains certifiably useful while effective support is sufficiently large relative to the observed gradient signal. Holding the other terms fixed, both the admissible step and the certified improvement shrink as ESS falls. This complements population policy-deviation guarantees, which characterize the consequences of changing the policy, by controlling the finite-rollout direction used to make that change. Having established when the IS update is reliable, we next ask when an alternative gradient estimator provides the stronger one-step certificate.

4.2 Choosing an alternative gradient estimator

The preceding result connects the reliability of the IS estimator to policy improvement. When this guarantee weakens, one may instead use a controlled coefficient rule, such as truncated IS or PPO-style masking. We compare a general rule u with IS in two steps: first as a question of gradient estimation, and then as a question of policy improvement at a common fixed step. Write $\hat{g}_k(u) := \hat{g}_N(u; \theta_k)$. For a rule u , adding and subtracting the estimator mean gives the MSE decomposition

$$\begin{aligned} m_k(u) &:= \mathbb{E} \|\hat{g}_k(u) - g_k\|_2^2 \\ &= \|\mathbb{E}[\hat{g}_k(u)] - g_k\|_2^2 + \mathbb{E} \|\hat{g}_k(u) - \mathbb{E}[\hat{g}_k(u)]\|_2^2. \end{aligned} \quad (15)$$

The two terms are squared bias and sampling variation. Since the IS estimator is unbiased, a controlled rule has lower MSE exactly when its reduction in sampling variation exceeds the squared bias introduced by changing the coefficients. We call this the estimator-level bias–variance crossover. The next theorem gives the expected-improvement certificate used to compare such estimators at a common fixed step.

Theorem 4 (Improvement certificate). *Suppose that J is L_k -smooth on every update segment generated by a coefficient rule u almost surely, for a constant $L_k > 0$, and that $\mathbb{E} \|\hat{g}_k(u)\|_2^2 < \infty$. For any step size $\eta_k > 0$,*

$$\mathbb{E} [J(\theta_k + \eta_k \hat{g}_k(u)) - J(\theta_k)] \geq \eta_k g_k^\top \mathbb{E}[\hat{g}_k(u)] - \frac{L_k \eta_k^2}{2} \mathbb{E} \|\hat{g}_k(u)\|_2^2. \quad (16)$$

Denote the right-hand side by

$$B_k(u; \eta_k) := \eta_k g_k^\top \mathbb{E}[\hat{g}_k(u)] - \frac{L_k \eta_k^2}{2} \mathbb{E} \|\hat{g}_k(u)\|_2^2. \quad (17)$$

The certificate has two parts. The first rewards alignment with the population gradient; the second penalizes the expected squared update under smoothness. Coefficient control improves the certificate only when the penalty it removes outweighs the useful alignment it removes. The next corollary states this crossover directly.

Corollary 5 (IS-to-controlled-estimator crossover). *Let u_{IS} denote the IS coefficient rule and let u be any alternative rule. Suppose the same deterministic L_k is valid almost surely for both families of update segments. With a common step size $0 < \eta_k \leq 1/L_k$, the alternative rule has the stronger certificate if and only if*

$$B_k(u; \eta_k) > B_k(u_{\text{IS}}; \eta_k) \iff m_k(u_{\text{IS}}) - m_k(u) > (1 - L_k \eta_k) \{ \mathbb{E} \|\hat{g}_k(u_{\text{IS}})\|_2^2 - \mathbb{E} \|\hat{g}_k(u)\|_2^2 \}. \quad (18)$$

When the alternative lowers the update second moment, the right-hand side of Equation (18) is its discounted reduction. Its MSE gain must exceed this amount. Since the IS estimator is unbiased, the same condition has the equivalent alignment form

$$\|g_k\|_2^2 - g_k^\top \mathbb{E}[\hat{g}_k(u)] < \frac{L_k \eta_k}{2} \{ \mathbb{E}\|\hat{g}_k(u_{\text{IS}})\|_2^2 - \mathbb{E}\|\hat{g}_k(u)\|_2^2 \}. \quad (19)$$

The left-hand side is the population-gradient alignment lost by changing the estimator; the right-hand side is the reduction in the quadratic smoothness penalty.

Remark 6 (Truncated importance sampling). Upper truncation, $u_{\text{T},t} = \min\{r_t, c\}$, caps each selected ratio without deleting its tail contribution. It is therefore a cheap variance safeguard with comparatively mild bias intervention. Its control is also limited: it acts only on the chosen upper tail, so the reduction in the update second moment can be modest under sequence-level mismatch.

Remark 7 (PPO-style masking). The PPO rule in Equation (6) removes saturated positive-advantage terms and low-ratio negative-advantage terms. It can reduce the update second moment more aggressively than upper truncation, but it also discards more population-gradient alignment. Equations (18) and (19) state when that stronger intervention is worthwhile.

5 Policy optimization and ESS

This section makes the theoretical distinction observable in two settings. An exactly solvable Gaussian bandit first separates gradient-estimation MSE from one-step policy improvement. Large-model trajectories then show how observed effective support evolves around delayed optimization failure.

5.1 Policy optimization in a Gaussian bandit

To isolate the estimator and optimization comparisons, we construct a scalar Gaussian bandit with a quadratic reward. Gaussian policies and smooth objectives provide a standard analytic setting for policy-gradient methods [5, 8, 10]. We hold the current policy fixed and vary only its distance from the rollout policy. Along this path, every quantity is available in closed form and contains no Monte Carlo error; Appendix A.1 gives the setup and calculations.

At each point on this path, the MSE oracle selects between IS and PPO by gradient-estimation MSE, whereas the certificate oracle selects the candidate with larger $B_k(u; \eta_k)$. The safe certificate oracle additionally includes a null update with certificate zero. To show how the alternative choices in Section 4.2 act, we also report upper-truncated IS with cap $c = 3$. We use $N = 32$, PPO radius $\epsilon = 0.2$, gradient $g = 2$, and step size $\eta = 0.4$, and vary ρ along one monotone ESS path.

Figure 2 first shows why MSE matters. As effective support contracts from $\rho = 0.961$ to 0.0109, the IS MSE rises from 0.228 to 62.3 and its exact expected improvement falls from 1.26 to -3.70 . The same ESS path therefore reveals the reliability loss before choosing a particular intervention. It is a regime diagnostic, not the oracle: the pairwise IS–PPO MSE crossover occurs at $\rho_{\text{MSE}} = 0.123$, while their certificate crossover occurs later at $\rho_B = 0.0392$.

The two controlled estimators explain the gap. TIS caps the upper tail while retaining substantially more population-gradient alignment. At the highest-support edge, TIS and IS are indistinguishable on the plotted scale; the substantive low-support crossover occurs at $\rho = 0.105$, below which TIS remains positive throughout the displayed path. PPO masking removes more of the update. It overtakes IS only below ρ_B and reaches zero certificate at $\rho_{\text{stop}} = 0.0109$, marked by the low-support reference line in Figure 2. These values are specific to $(g, N, \epsilon, \eta, c)$, but the

ordering is the point: ESS locates the support regime, while MSE reduction and retained alignment jointly determine the better update.




Figure 2: Exact Gaussian-bandit comparison with $g = 2$, $N = 32$, $\epsilon = 0.2$, $c = 3$, and $\eta = 0.4$. (a) Gradient MSE as effective support contracts. (b) Population-gradient alignment retained by each estimator. (c) Exact expected one-step improvement, equal here to $B_u(\eta)$. The substantive IS–TIS crossover is $\rho = 0.105$; the dashed line marks the IS–PPO crossover at $\rho_B = 0.0392$, and the dash-dotted line marks PPO’s zero certificate at $\rho_{\text{stop}} = 0.0109$.

5.2 Policy optimization in large language models

The Gaussian comparison holds the policy state fixed. Large-model training adds a temporal question: does observed effective support deteriorate before the performance loss becomes visible? Figure 3 places the two permissive Qwen3-30B-A3B trajectories from Figure 1, which use TIS caps 3 and 5, beside a third trajectory trained with GRPO without clipping. Each run is trained for up to 200 policy updates and evaluated on AIME-2024 using mean@16.

The two TIS trajectories establish the temporal ordering that motivates an ESS diagnostic. Their sample ESS first falls below 0.01 at steps 77 and 68, respectively, after both performance curves peak at step 60 and before their subsequent collapse. The unclipped trajectory supplies the necessary counterexample: its first sub-0.01 observation at step 57 is isolated, ESS recovers at the next updates, and AIME performance continues from 28.1% at step 60 to a 44.4% peak at step 140. Low ESS therefore identifies weak effective support without itself certifying immediate failure or determining which controlled estimator gives the better update.




Figure 3: Effective support along three permissive Qwen3-30B-A3B trajectories. Panels (a) and (b) reproduce the TIS-3 and TIS-5 motivation runs, whose first crossings below the 0.01 ESS reference precede collapse. Panel (c) shows GRPO without clipping: its crossing at step 57 is transient, and performance continues improving to step 140. The black curves report AIME-2024 mean@16 and the red curves report per-step normalized sequence ESS. Vertical dotted lines mark the first sub-0.01 observation.

The two examples give ESS a precise empirical role. In the Gaussian bandit, ESS indexes the support regime, while the full certificate determines the preferred update. In the language-model trajectories, sample ESS reveals support loss before failure becomes visible, but does not itself constitute a failure certificate. The operational question is therefore whether conditioning a safeguard on ESS improves low-support updates without sacrificing permissive learning when support is high. We test this rule next.

6 ESS-conditioned policy optimization

Section 5 identifies ESS as a coordinate for effective support. A natural extension is to test whether this coordinate can time an existing safeguard: retain a permissive update while observed support is broad, and activate control only after support contracts. We first study this transition in Qwen3-30B-A3B math RLVR, then change the model scale and training objective.

6.1 Update rule and protocol

Let $\hat{g}_{\text{perm}}$ denote the permissive update and $\hat{g}_{\text{safe}}$ a prespecified safeguard. Given a threshold τ , the ESS-conditioned update is

$$\hat{g}_{\text{gate}} = \begin{cases} \hat{g}_{\text{perm}}, & \hat{\rho} \geq \tau, \\ \hat{g}_{\text{safe}}, & \hat{\rho} < \tau. \end{cases} \quad (20)$$

ESS determines when the update changes; the selected safeguard determines how the low-support update is controlled. We fix $\tau = 0.1$ throughout. This value was prescribed for the language-model runs and is not derived from the Gaussian thresholds in Section 5.1.

The main experiment trains Qwen3-30B-A3B for 200 policy updates on the boxed mathematical-reasoning data of Yu et al. [18] and evaluates AIME-2024 using mean@16. The permissive branch uses truncated importance sampling (TIS) with ratio cap 3, or no coefficient clipping in the explicit

unclipped ablation. The low-support branch uses conventional GRPO clipping, Clip-Higher [18], DPPO [11], or skips the minibatch update. Most variants compute $\hat{\rho}$ from the raw complete-sequence ratios before coefficient control. The 30B TIS/Clip-Higher and TIS-5 constructions, their 8B counterpart, and the RLHF transfer instead compute the same statistic after the coefficient cap; we refer to this implementation as shaped ESS. Results are complete training trajectories, so the comparisons below describe the temporal and endpoint behavior of the reported runs. Each curve represents one training trajectory; no across-seed uncertainty is available.

6.2 Main results

Figure 4 compares four 30B constructions that pair a permissive high-support update with a prespecified low-support safeguard. Panels (a) and (b) are complete-recipe comparisons matched on the safeguard’s clip band; panels (c) and (d) directly compare static and ESS-conditioned versions of the named estimator.




Figure 4: ESS-conditioned policy optimization on Qwen3-30B-A3B. (a) Static GRPO versus TIS 3 with an ESS-conditioned GRPO safeguard. (b) Static Clip-Higher versus TIS 3 with an ESS-conditioned Clip-Higher safeguard. (c) Static versus ESS-conditioned DPPO. (d) Static TIS 3 versus the same TIS-3/Clip-Higher trajectory shown in panel (b). Direct labels give the last reported AIME-2024 mean@16; static TIS 3 ends at step 170 after collapse, while the other trajectories end at step 200.

GRPO. Static GRPO peaks at 29.6% and finishes at 26.2%. The conditional recipe, which uses TIS 3 while ESS is high and the standard 0.2/0.2 GRPO band when ESS is low, peaks at 40.8%

and finishes at 38.5%. Because the high-support branches differ, this is a band-matched comparison of complete recipes rather than a gate-only ablation.

Clip-Higher. Always applying the 0.2/0.28 Clip-Higher band reaches and finishes at 37.9%. Using TIS 3 in the high-support regime and activating the same band only in the low-support regime instead peaks at 47.3% and finishes at 44.8%. This is likewise a complete-recipe comparison, now holding the low-support safeguard fixed.

DPPO. Static DPPO peaks at 37.5% but recedes to 30.4%, whereas its ESS-conditioned counterpart peaks at 40.2% and finishes at 39.2%. Conditioning therefore preserves the late trajectory without sacrificing its best observed performance.

TIS 3. Without an ESS gate, TIS 3 reaches 31.7% at step 60 and then collapses to 0.417% at its last recorded evaluation, step 170. The conditional TIS-3/Clip-Higher construction instead reaches 47.3% and finishes at 44.8%. Its trajectory exhibits the intended ordering: permissive learning first, followed by increasingly controlled updates as effective support contracts.

Unclipped update. The unclipped comparison isolates the intervention rule more directly. All three runs in Figure 5 use the same unclipped coefficient update. The reference has no ESS-conditioned action; the other two runs either clip or skip when raw ESS is below 0.1. The reference attains the highest early peak, 44.4%, but finishes at 34.6%. Conditional clipping retains nearly the same peak, 43.3%, and improves the endpoint to 38.1%. Skipping the low-support updates finishes at 35.6%, but reaches that level more slowly. Thus, in this controlled family, clipping preserves most of the permissive learning while providing the strongest late retention.




Figure 5: Gate-action ablation with an unclipped Qwen3-30B-A3B update. The reference run has no ESS-conditioned action; its nominal fraction trigger produces only three small interventions. The other runs either clip or skip a minibatch update when raw normalized ESS is below 0.1. Direct labels give the step-200 AIME-2024 mean@16.

The implemented intervention also follows the predicted temporal ordering. For the DPPO construction, mean raw ESS falls from 0.412 over steps 1–50 to 0.120 over steps 151–200, while the mean intervention fraction rises from 0.153 to 0.668. Figure 9 reports the corresponding patterns for all three safeguards. Protection is therefore present but limited early, and intensifies as effective support contracts.

TIS 5. The TIS-5 pair isolates the gate more directly because both trajectories use the same permissive update and differ only in the ESS-conditioned skip. Static TIS 5 peaks at 28.3% at step 60 and then collapses to 0.00%, whereas the conditioned trajectory reaches and finishes at 39.4%. This replication therefore shows that the late rescue is not specific to the TIS-3 cap or to clipping as the low-support action.




Figure 6: Gate-only replication with TIS cap 5. Panel (a) compares the AIME-2024 mean@16 trajectories; panel (b) reports normalized shaped ESS. The static trajectory collapses to 0.00%, whereas the fixed-0.1 ESS-conditioned skip reaches 39.4%.

6.3 Transfer across models and objectives

We first change the model scale while keeping the math-RLVR task fixed. Figure 7 reports the two Qwen3-8B comparisons for which loss aggregation is matched. Under Clip-Higher, the static trajectory finishes at 29.0%, while the ESS-conditioned construction finishes at 32.7%. Under DPPO, the corresponding endpoints are 31.7% and 32.5%. These gains are smaller than the 30B rescue and depend on the safeguard. The endpoint ordering favors the conditional construction in both aggregation-matched comparisons after the change in model scale.




Figure 7: Dense-model transfer on Qwen3-8B. Panel (a) compares token-mean Clip-Higher with its token-mean ESS-conditioned construction. Panel (b) compares static and ESS-conditioned DPPO under the same sum-normalized aggregation. Direct labels give the step-200 AIME-2024 mean@16.

Finally, we change both the objective and the data. We train Qwen3-4B-Instruct-2507 for 200 updates on Anthropic HH-RLHF [1] using Skywork-Reward-Llama-3.1-8B [4] as the reward model. Figure 8 compares the complete GRPO and TIS-3-plus-ESS recipes. Their reward-model scores start at comparable levels and finish at 43.0 and 60.7, respectively. This experiment tests optimization transfer on the training reward model; it does not measure human preference directly.




Figure 8: Transfer to open-ended RLHF on Qwen3-4B-Instruct-2507. Thin curves are per-step reward-model scores; thick curves are trailing seven-step means. Direct labels give the step-200 observations. This is a complete-recipe comparison between GRPO and TIS 3 with an ESS-conditioned safeguard.

Across these experiments, ESS contributes timing rather than a new clipping rule. High observed support leaves the permissive branch active; lower support increases the use of a prespecified control. The 30B rescue patterns, the smaller 8B gains, and the RLHF transfer are consistent with the same division of labor: ESS locates the reliability regime, while the safeguarded estimator determines the low-support update.

7 Conclusion

We studied clipping as a finite-sample estimator choice rather than a default property of policy optimization. The theory separates gradient-estimation error from the useful update retained after control: reducing MSE improves the fixed-step certificate only when that gain compensates for the signal removed by clipping. This distinction makes IS preferable while effective support is broad and explains why controlled estimators become attractive as mismatch grows. The exact Gaussian construction exhibits separate MSE and update crossovers. The reported language-model trajectories then show that sequence ESS can reveal support contraction and time a prespecified safeguard without serving as the oracle itself. In the RLVR and RLHF examples, the resulting conditional constructions retain permissive learning in the higher-support regime and protect training more strongly as support contracts. ESS therefore provides the operational link between importance-weight reliability and the decision of when to control an update.

A Numerical Experiments

A.1 Gaussian bandit

This section gives the calculations used in Section 5.1.

Setup We compare IS, upper-truncated IS with $c = 3$, and the PPO-style rule in Equation (6), using the exact current-policy value baseline.

Because the score contribution is polynomial and each truncation or mask boundary is linear in z , the means, MSEs, and update second moments are finite sums of truncated Gaussian moments.

We choose a quadratic reward because it makes the smoothness certificate equal to the true expected gain. Any separation between the two oracles is therefore genuine rather than an artifact of a loose bound.

Consider the current policy and rollout policy

$$P_\theta = \mathcal{N}(\theta, 1), \quad Q_\Delta = \mathcal{N}(\theta - \Delta, 1), \quad \Delta \geq 0, \quad (21)$$

with reward $R(a) = -\frac{1}{2}(a - a^\star)^2$. We hold the current policy fixed and vary only the rollout gap Δ . Writing $g := a^\star - \theta > 0$, the objective is

$$J(\theta) = -\frac{1}{2}\{(\theta - a^\star)^2 + 1\}, \quad \nabla J(\theta) = g, \quad L = 1. \quad (22)$$

Here θ is the fixed pre-update mean and Q_Δ remains frozen during the hypothetical one-step update. We vary Δ only across common-state comparisons. For $z = a - \theta$, the likelihood ratio and normalized ESS reduce to

$$W_\Delta(z) = \exp\{\Delta z + \Delta^2/2\}, \quad \rho(\Delta) = \{\mathbb{E}_{Q_\Delta}[W_\Delta^2]\}^{-1} = \exp(-\Delta^2). \quad (23)$$

Normalized ESS is therefore a one-to-one coordinate for this branch.

Write $\mu_u := \mathbb{E}[\hat{g}_u]$, $s_u := \mathbb{E}[\hat{g}_u^2]$, and $B_u(\eta) := \eta g \mu_u - \eta^2 s_u/2$. For this quadratic objective, the fixed-step certificate is exact:

$$\mathbb{E}[J(\theta + \eta \hat{g}_u) - J(\theta)] = \eta g \mu_u - \frac{\eta^2}{2} s_u = B_u(\eta). \quad (24)$$

We use $N = 32$, the standard PPO radius $\epsilon = 0.2$, TIS cap $c = 3$, and choose $g = 2$ and $\eta = 0.4$ so that the MSE, improvement, and abstention regimes are all visible on one monotone ESS path.

Exact Closed Form Computations Let $z = a - \theta$. Under the current policy, $z \sim \mathcal{N}(0, 1)$ and the score is z . The exact current-policy value baseline gives

$$A_g(z) = gz + \frac{1 - z^2}{2}, \quad f_g(z) := zA_g(z) = gz^2 + \frac{z - z^3}{2}. \quad (25)$$

Although f_g is unbounded, all polynomial Gaussian moments used below are finite. The example therefore uses the finite-moment fixed-step certificate; the bounded-integrand IS theorem addresses a separate reliability question. IS uses $W_\Delta f_g$. The PPO-style candidate uses $W_\Delta f_g M_{\epsilon, \Delta}$, where

$$M_{\epsilon, \Delta}(z) := \mathbf{1}\{A_g(z) \geq 0, W_\Delta(z) \leq 1 + \epsilon\} + \mathbf{1}\{A_g(z) < 0, W_\Delta(z) \geq 1 - \epsilon\}. \quad (26)$$

Let $M_{\text{IS}} \equiv 1$ and $M_{\text{P}} = M_{\epsilon, \Delta}$. Change of measure gives the exact per-sample moments

$$\begin{aligned}\mu_u(\Delta) &= \mathbb{E}_{U \sim \mathcal{N}(0,1)}[f_g(U)M_u(U)], \\ h_u(\Delta) &= e^{\Delta^2} \mathbb{E}_{V \sim \mathcal{N}(\Delta,1)}[f_g(V)^2 M_u(V)], \quad u \in \{\text{IS}, \text{P}\}.\end{aligned}\tag{27}$$

For TIS, which in this one-factor bandit uses $\min\{W_\Delta, c\}f_g$, let $b_c(\Delta) := \{\log c - \Delta^2/2\}/\Delta$. Since $\min\{W_\Delta, c\} = W_\Delta$ below $b_c(\Delta)$ and equals c above it,

$$\begin{aligned}\mu_{\text{T}}(\Delta) &= \int_{-\infty}^{b_c(\Delta)} f_g(z)\phi(z) dz + c \int_{b_c(\Delta)}^{\infty} f_g(z)\phi(z + \Delta) dz, \\ h_{\text{T}}(\Delta) &= e^{\Delta^2} \int_{-\infty}^{b_c(\Delta)} f_g(z)^2 \phi(z - \Delta) dz + c^2 \int_{b_c(\Delta)}^{\infty} f_g(z)^2 \phi(z + \Delta) dz.\end{aligned}\tag{28}$$

In particular, $\mu_{\text{IS}} = g$. For an iid estimator with batch size N ,

$$m_u = (\mu_u - g)^2 + \frac{h_u - \mu_u^2}{N}, \quad s_u = \mu_u^2 + \frac{h_u - \mu_u^2}{N}.\tag{29}$$

The advantage changes sign at $g \pm \sqrt{g^2 + 1}$, and the ratio boundaries are $(\log(1 \pm \epsilon) - \Delta^2/2)/\Delta$. Equation (27) is therefore a finite sum of truncated normal polynomial moments. With $I_j(a, b) := \int_a^b z^j \phi(z) dz$, integration by parts gives

$$I_j(a, b) = a^{j-1}\phi(a) - b^{j-1}\phi(b) + (j-1)I_{j-2}(a, b), \quad j \geq 2,\tag{30}$$

starting from $I_0 = \Phi(b) - \Phi(a)$ and $I_1 = \phi(a) - \phi(b)$. Boundary terms vanish at infinity. Substituting these moments into Equation (29) and the exact gain in Equation (24) gives the reported boundaries.

Results We evaluate $0.2 \leq \Delta \leq \sqrt{-\log(0.005)}$. Numerical evaluation of the closed-form moments gives four substantive boundaries on this branch:

$$\rho_{\text{MSE}} = 0.123, \quad \rho_{\text{T}} = 0.105, \quad \rho_B = 0.0392, \quad \rho_{\text{stop}} = 0.0109.\tag{31}$$

The MSE oracle over $\mathcal{U} = \{u_{\text{IS}}, u_{\text{P}}\}$ selects IS above ρ_{MSE} and PPO below it. The safe certificate oracle adds the null update and instead selects IS for $\rho > \rho_B$, PPO for $\rho_{\text{stop}} < \rho < \rho_B$, and no update for $\rho < \rho_{\text{stop}}$. Adjacent choices tie at each boundary. The TIS overlay has lower MSE than both pairwise candidates on the displayed branch; its certificate overtakes IS at ρ_{T} and remains above the PPO certificate thereafter. Because the cap barely changes the upper tail at the highest-support edge, IS and TIS also cross at $\rho = 0.922$ with a maximum certificate difference of only 2.64×10^{-6} over $0.922 < \rho \leq 0.961$. They are visually indistinguishable there; $\rho_{\text{T}} = 0.105$ denotes the substantive low-support crossover discussed in the main text.

B Additional coefficient-rule details

The main text treats coefficient control generically because different rules alter different parts of the gradient. A factorwise detached upper truncation uses $u_{\text{T},t}(\mathbf{r}, A) := \min\{r_t, c\}$. Its contribution obeys

$$\|\psi_k(u_{\text{T}}; Z)\|_2 \leq c|A(Z)| \sum_{t=1}^T \|\ell_{\theta_k,t}(Z)\|_2,\tag{32}$$

so it has a direct second-moment envelope whenever the right-hand side has a finite second moment.

PPO masking is different. For $A \geq 0$ and $c = 1 + \epsilon$,

$$0 \leq r - \min\{r, c\} = (r - c)_+ \leq r \mathbf{1}\{r > c\} = r - rM_\epsilon(r, A), \quad (33)$$

where $M_\epsilon(r, A) = \mathbf{1}\{A \geq 0, r \leq 1 + \epsilon\} + \mathbf{1}\{A < 0, r \geq 1 - \epsilon\}$. Truncation retains a bounded positive coefficient above the cap, whereas PPO removes that contribution. On the negative-advantage lower branch, PPO masks small ratios whereas upper truncation does not. A complete-sample mask is another distinct rule: it sets $u_t(\mathbf{r}, A) = W_k M_\epsilon(W_k, A)$ for every factor. We state which rule is used in each experiment rather than treating these implementations as interchangeable.

C Additional RLVR diagnostics

C.1 Intervention and training dynamics

The main text reports evaluation performance. We record here how often the conditional rule intervenes and the accompanying training dynamics.




Figure 9: Realized ESS-conditioned intervention schedules. Each panel pairs the per-step intervention fraction (thin black) and its trailing seven-step mean (thick black) with normalized ESS (red). GRPO and DPPO use raw ESS; Clip-Higher uses shaped ESS after the coefficient cap. The dotted line marks the threshold 0.1.

Figure 9 shows the intended temporal ordering: intervention is already possible early, but its frequency increases as ESS contracts. The remaining figures record response length, entropy, and reward diagnostics.




Figure 10: Training dynamics for the main comparisons. The static TIS 3 run develops extreme response length and entropy together with a collapse in training reward. The safeguards and their ESS-conditional variants avoid this failure over the reported 200 steps.




Figure 11: Detailed TIS 3 comparison. Under static TIS, ESS falls below 0.01 after step 77 and the response length and entropy diverge. The fixed-0.1 intervention increases the clipped fraction as ESS declines and keeps both diagnostics controlled.

D Proofs for the theoretical results

At update k , all expectations in this appendix are taken after conditioning on the preceding updates. The current policy is therefore fixed and the unused sampling units are independent draws from Q . We suppress this conditioning throughout.

D.1 Policy-gradient identity and mean-square bound

By definition, $f_k = (R - b(X))\nabla_\theta \log P_{\theta_k}(Y | X)$ and $W_k = dP_{\theta_k}/dQ$. Change of measure gives

$$\begin{aligned}\mathbb{E}_Q[W_k f_k] &= \mathbb{E}_{P_{\theta_k}}[(R - b(X))\nabla_\theta \log P_{\theta_k}(Y | X)] \\ &= \mathbb{E}_{P_{\theta_k}}[R\nabla_\theta \log P_{\theta_k}(Y | X)] \\ &\quad - \mathbb{E}_{P_{\theta_k}}[b(X)\nabla_\theta \log P_{\theta_k}(Y | X)].\end{aligned}\tag{34}$$

For the baseline term, condition further on X . Under the standard regularity condition that permits differentiation under the conditional integral,

$$\begin{aligned}\mathbb{E}_{P_{\theta_k}}[b(X)\nabla_\theta \log P_{\theta_k}(Y | X) | X] &= b(X) \int P_{\theta_k}(y | X) \nabla_\theta \log P_{\theta_k}(y | X) d\lambda(y) \\ &= b(X) \nabla_\theta \int P_{\theta_k}(y | X) d\lambda(y) \\ &= b(X) \nabla_\theta 1 = 0.\end{aligned}\tag{35}$$

The remaining term is the score-function identity. Hence

$$\mathbb{E}_Q[W_k f_k] = \nabla J(\theta_k) = g_k.\tag{36}$$

Averaging N independent copies gives $\mathbb{E}[\hat{g}_k] = g_k$.

For the mean-square calculation, define $\xi_{k,i} := W_{k,i} f_{k,i} - g_k$. Equation (36) implies $\mathbb{E}[\xi_{k,i}] = 0$. Since $\hat{g}_k - g_k = N^{-1} \sum_{i=1}^N \xi_{k,i}$,

$$\mathbb{E}\|\hat{g}_k - g_k\|_2^2 = \frac{1}{N^2} \sum_{i=1}^N \mathbb{E}\|\xi_{k,i}\|_2^2 + \frac{1}{N^2} \sum_{i \neq j} \mathbb{E}[\xi_{k,i}^\top \xi_{k,j}].\tag{37}$$

For $i \neq j$, independence and zero means give

$$\mathbb{E}[\xi_{k,i}^\top \xi_{k,j}] = \mathbb{E}[\xi_{k,i}]^\top \mathbb{E}[\xi_{k,j}] = 0.\tag{38}$$

The diagonal terms are identical, so

$$\mathbb{E}\|\hat{g}_k - g_k\|_2^2 = \frac{1}{N} \mathbb{E}_Q \|W_k f_k - g_k\|_2^2.\tag{39}$$

Expanding the remaining squared norm yields

$$\begin{aligned}\mathbb{E}_Q \|W_k f_k - g_k\|_2^2 &= \mathbb{E}_Q [W_k^2 \|f_k\|_2^2] - 2g_k^\top \mathbb{E}_Q [W_k f_k] + \|g_k\|_2^2 \\ &= \mathbb{E}_Q [W_k^2 \|f_k\|_2^2] - \|g_k\|_2^2,\end{aligned}\tag{40}$$

where the last equality uses Equation (36). Combining Equations (39) and (40) gives

$$\mathbb{E}\|\widehat{g}_k - g_k\|_2^2 = \frac{1}{N} \left\{ \mathbb{E}_Q[W_k^2 \|f_k\|_2^2] - \|g_k\|_2^2 \right\}. \quad (41)$$

Finally, $\|f_k(z)\|_2 \leq \|f_k\|_\infty$ and $\mathbb{E}_Q[W_k^2] = \rho_k^{-1}$ imply

$$\begin{aligned} \mathbb{E}\|\widehat{g}_k - g_k\|_2^2 &\leq \frac{1}{N} \mathbb{E}_Q[W_k^2 \|f_k\|_2^2] \\ &\leq \frac{\|f_k\|_\infty^2}{N} \mathbb{E}_Q[W_k^2] \\ &= \frac{\|f_k\|_\infty^2}{N \rho_k}, \end{aligned} \quad (42)$$

which proves Equation (11).

D.2 Proof of Lemma 2

Apply Markov's inequality to the nonnegative random variable $\|\widehat{g}_k - g_k\|_2^2$:

$$\begin{aligned} \Pr \left(\|\widehat{g}_k - g_k\|_2 > \frac{\|f_k\|_\infty}{\sqrt{\delta N \rho_k}} \right) \\ &= \Pr \left(\|\widehat{g}_k - g_k\|_2^2 > \frac{\|f_k\|_\infty^2}{\delta N \rho_k} \right) \\ &\leq \frac{\mathbb{E}\|\widehat{g}_k - g_k\|_2^2}{\|f_k\|_\infty^2 / (\delta N \rho_k)} \\ &\leq \delta, \end{aligned} \quad (43)$$

where the last line uses Equation (11). Taking complements proves Equation (12).

D.3 Proof of Lemma 1

Let $\|\widehat{g} - g\|_2 \leq \varepsilon$. Smoothness gives, for every $\gamma \geq 0$,

$$J(\theta + \gamma \widehat{g}) - J(\theta) \geq \gamma g^\top \widehat{g} - \frac{L\gamma^2}{2} \|\widehat{g}\|_2^2. \quad (44)$$

The inner product satisfies

$$\begin{aligned} g^\top \widehat{g} &= \|\widehat{g}\|_2^2 + (g - \widehat{g})^\top \widehat{g} \\ &\geq \|\widehat{g}\|_2^2 - \|g - \widehat{g}\|_2 \|\widehat{g}\|_2 \\ &\geq \|\widehat{g}\|_2 (\|\widehat{g}\|_2 - \varepsilon). \end{aligned} \quad (45)$$

Substituting Equation (45) into Equation (44) gives

$$J(\theta + \gamma \widehat{g}) - J(\theta) \geq \gamma \|\widehat{g}\|_2 (\|\widehat{g}\|_2 - \varepsilon) - \frac{L\gamma^2}{2} \|\widehat{g}\|_2^2. \quad (46)$$

For $\widehat{g} \neq 0$, the derivative of the right-hand side is

$$\|\widehat{g}\|_2 (\|\widehat{g}\|_2 - \varepsilon) - L\gamma \|\widehat{g}\|_2^2. \quad (47)$$

Its nonnegative maximizer is $\gamma = L^{-1}(1 - \varepsilon/\|\widehat{g}\|_2)_+$. When $\widehat{g} = 0$, both the chosen step and the claimed lower bound are zero. Substituting the maximizer into Equation (46) proves Equation (9).

D.4 Proof of Theorem 3

Lemma 2 shows that the event

$$\|\widehat{g}_k - g_k\|_2 \leq \frac{\|f_k\|_\infty}{\sqrt{\delta N \rho_k}} \quad (48)$$

has probability at least $1 - \delta$. On this event, apply Lemma 1 with $\theta = \theta_k$, $\widehat{g} = \widehat{g}_k$, $\varepsilon = \|f_k\|_\infty / \sqrt{\delta N \rho_k}$, and $L = L_k$. The resulting inequality is exactly Equation (14).

D.5 Risk and certificate comparisons for coefficient rules

For a coefficient rule u , add and subtract the estimator mean:

$$\widehat{g}_k(u) - g_k = \{\widehat{g}_k(u) - \mathbb{E}[\widehat{g}_k(u)]\} + \{\mathbb{E}[\widehat{g}_k(u)] - g_k\}. \quad (49)$$

Squaring and taking expectation gives

$$\begin{aligned} m_k(u) &= \mathbb{E} \|\widehat{g}_k(u) - \mathbb{E}[\widehat{g}_k(u)]\|_2^2 + \|\mathbb{E}[\widehat{g}_k(u)] - g_k\|_2^2 \\ &\quad + 2\{\mathbb{E}[\widehat{g}_k(u)] - g_k\}^\top \mathbb{E}[\widehat{g}_k(u) - \mathbb{E}[\widehat{g}_k(u)]] \\ &= \mathbb{E} \|\widehat{g}_k(u) - \mathbb{E}[\widehat{g}_k(u)]\|_2^2 + \|\mathbb{E}[\widehat{g}_k(u)] - g_k\|_2^2, \end{aligned} \quad (50)$$

because the centered term has mean zero. This proves Equation (15). For an iid sample mean, the first term in Equation (50) is the per-sample variation divided by N .

It remains to connect total risk to the expected policy-improvement certificate. For any random estimate $\widehat{g}$ of g_k , smoothness gives

$$J(\theta_k + \eta_k \widehat{g}) - J(\theta_k) \geq \eta_k g_k^\top \widehat{g} - \frac{L_k \eta_k^2}{2} \|\widehat{g}\|_2^2. \quad (51)$$

Taking expectation in Equation (51), with $\mathbb{E}\|\widehat{g}\|_2^2 < \infty$, gives

$$\mathbb{E}[J(\theta_k + \eta_k \widehat{g}) - J(\theta_k)] \geq \eta_k g_k^\top \mathbb{E}[\widehat{g}] - \frac{L_k \eta_k^2}{2} \mathbb{E}\|\widehat{g}\|_2^2. \quad (52)$$

Setting $\widehat{g} = \widehat{g}_k(u)$ proves Theorem 4. The polarization identity, followed by expectation, gives

$$2g_k^\top \mathbb{E}[\widehat{g}] = \|g_k\|_2^2 + \mathbb{E}\|\widehat{g}\|_2^2 - \mathbb{E}\|\widehat{g} - g_k\|_2^2. \quad (53)$$

Substituting Equation (53) into the certificate and subtracting the IS certificate gives

$$\begin{aligned} B_k(u; \eta_k) - B_k(u_{\text{IS}}; \eta_k) &= \frac{\eta_k}{2} [m_k(u_{\text{IS}}) - m_k(u)] \\ &\quad - \frac{\eta_k}{2} (1 - L_k \eta_k) \{\mathbb{E}\|\widehat{g}_k(u_{\text{IS}})\|_2^2 - \mathbb{E}\|\widehat{g}_k(u)\|_2^2\}. \end{aligned} \quad (54)$$

Since $\eta_k > 0$, the right-hand side is positive exactly under Equation (18), proving Corollary 5. Since the IS estimator is unbiased, subtracting the certificate definitions instead gives

$$\begin{aligned} B_k(u; \eta_k) - B_k(u_{\text{IS}}; \eta_k) &= -\eta_k \left\{ \|g_k\|_2^2 - g_k^\top \mathbb{E}[\widehat{g}_k(u)] \right\} \\ &\quad + \frac{L_k \eta_k^2}{2} \{\mathbb{E}\|\widehat{g}_k(u_{\text{IS}})\|_2^2 - \mathbb{E}\|\widehat{g}_k(u)\|_2^2\}. \end{aligned} \quad (55)$$

Its positivity is equivalent to Equation (19).

D.6 Intact prompt groups

Suppose M prompts are independent and each prompt has G conditionally independent responses. Keep each prompt group intact and define its contribution by $\Xi_{k,i} := G^{-1} \sum_{j=1}^G W_{k,ij} f_{k,ij}$. A leave-one-out baseline is admissible because it excludes the focal response. Conditional on the prompt and the remaining responses, the focal score still has mean zero. Hence $\mathbb{E}[\Xi_{k,i}] = g_k$.

Jensen’s inequality gives

$$\begin{aligned} \|\Xi_{k,i}\|_2^2 &= \left\| \frac{1}{G} \sum_{j=1}^G W_{k,ij} f_{k,ij} \right\|_2^2 \\ &\leq \frac{1}{G} \sum_{j=1}^G W_{k,ij}^2 \|f_{k,ij}\|_2^2. \end{aligned} \quad (56)$$

Repeating the sample-mean calculation across the M independent prompt groups gives

$$\begin{aligned} \mathbb{E} \left\| \frac{1}{M} \sum_{i=1}^M \Xi_{k,i} - g_k \right\|_2^2 &\leq \frac{1}{M} \mathbb{E} \|\Xi_{k,1}\|_2^2 \\ &\leq \frac{\|f_k\|_\infty^2}{M \rho_k}. \end{aligned} \quad (57)$$

Markov’s inequality then gives the same form as Lemma 2 with the number of independent prompts M in place of the number of responses N .

D.7 Finite-moment relaxation

The exact equality in Equation (10) requires only $\mathbb{E}_Q[W_k^2 \|f_k\|_2^2] < \infty$. The bounded-contribution condition is used solely to obtain the standard form $\|f_k\|_\infty^2 / (N \rho_k)$. More generally, define $M_{2,k} := \mathbb{E}_Q[W_k^2 \|f_k\|_2^2] / \mathbb{E}_Q[W_k^2]$. The same calculation gives $\mathbb{E} \|\hat{g}_k - g_k\|_2^2 \leq M_{2,k} / (N \rho_k)$. We use the bounded form in the main text to match conventional importance-sampling notation and make the role of normalized ESS explicit.

References

- [1] Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, Nicholas Joseph, Saurav Kadavath, Jackson Kernion, Tom Conerly, et al. Training a helpful and harmless assistant with reinforcement learning from human feedback, 2022.
- [2] Rasool Fakoor, Murdock Aubry, Nicholas Stranges, and Alexander J. Smola. Trust the batch, on- or off-policy: Adaptive policy optimization for RL post-training, 2026.
- [3] Luke J. Huang, Zhuoyang Zhang, Qinghao Hu, Shang Yang, and Song Han. Stable asynchrony: Variance-controlled off-policy RL for LLMs, 2026.
- [4] Chris Yuhao Liu, Liang Zeng, Jiakai Liu, Rui Yan, Jujie He, Chaojie Wang, Shuicheng Yan, Yang Liu, and Yahui Zhou. Skywork-Reward: Bag of tricks for reward modeling in LLMs, 2024.

- [5] Alberto Maria Metelli, Matteo Papini, Francesco Faccio, and Marcello Restelli. Policy optimization via importance sampling. In *Advances in Neural Information Processing Systems*, volume 31, 2018. URL <https://arxiv.org/abs/1809.06098>.
- [6] Alberto Maria Metelli, Matteo Papini, Nico Montali, and Marcello Restelli. Importance sampling techniques for policy optimization. *Journal of Machine Learning Research*, 21(141): 1–75, 2020. URL <https://jmlr.org/papers/v21/20-124.html>.
- [7] MiniMax. MiniMax-M1: Scaling test-time compute efficiently with lightning attention, 2025.
- [8] Matteo Papini, Matteo Pirotta, and Marcello Restelli. Smoothing policies and safe policy gradients. *Machine Learning*, 111:4081–4137, 2022. doi: 10.1007/s10994-022-06232-6. URL <https://doi.org/10.1007/s10994-022-06232-6>.
- [9] Jaesung R. Park, Junsu Kim, Gyeongman Kim, Jinyoung Jo, Sean Choi, Jaewoong Cho, and Ernest K. Ryu. Clip-low increases entropy and clip-high decreases entropy in reinforcement learning of large language models, 2025.
- [10] Matteo Pirotta, Marcello Restelli, and Luca Bascetta. Adaptive step-size for policy gradient methods. In *Advances in Neural Information Processing Systems*, volume 26, 2013. URL https://proceedings.neurips.cc/paper_files/paper/2013/hash/f64eac11f2cd8f0efa196f8ad173178e-Abstract.html.
- [11] Penghui Qi, Xiangxin Zhou, Zichen Liu, Tianyu Pang, Chao Du, Min Lin, and Wee Sun Lee. Rethinking the trust region in LLM reinforcement learning, 2026.
- [12] Alfréd Rényi. On measures of entropy and information. In *Proceedings of the Fourth Berkeley Symposium on Mathematical Statistics and Probability*, volume 1, pages 547–561. University of California Press, 1961. URL <https://projecteuclid.org/euclid.bsmsp/1200512181>.
- [13] John Schulman, Sergey Levine, Philipp Moritz, Michael I. Jordan, and Pieter Abbeel. Trust region policy optimization. In *Proceedings of the 32nd International Conference on Machine Learning*, volume 37 of *Proceedings of Machine Learning Research*, pages 1889–1897. PMLR, 2015. URL <https://proceedings.mlr.press/v37/schulman15.html>.
- [14] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms, 2017.
- [15] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models, 2024.
- [16] Guobin Shen, Chenxiao Zhao, Xiang Cheng, Lei Huang, and Xing Yu. VESPO: Variational sequence-level soft policy optimization for stable off-policy LLM training, 2026.
- [17] Tim van Erven and Peter Harremoës. Rényi divergence and Kullback–Leibler divergence. *IEEE Transactions on Information Theory*, 60(7):3797–3820, 2014. doi: 10.1109/TIT.2014.2320500.
- [18] Qiyang Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming Sheng, Yuxuan Tong, Chi Zhang, Mofan Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu, Jiaze Chen, Jiangjie Chen, Chengyi Wang, Hongli Yu, Yuxuan Song, Xiangpeng Wei, Hao Zhou,

Jingjing Liu, Wei-Ying Ma, Ya-Qin Zhang, Lin Yan, Mu Qiao, Yonghui Wu, and Mingxuan Wang. DAPO: An open-source LLM reinforcement learning system at scale, 2025.

- [19] Haizhong Zheng, Jiawei Zhao, and Beidi Chen. Prosperity before collapse: How far can off-policy RL reach with stale data on LLMs?, 2025.