

Rethinking the policy optimization with Gradient Estimators

Dongkyu Derek Cho, Thomas Voice

Abstract

Permissive policy optimization can learn rapidly from rollout data and then fail abruptly after many successful updates. Policy-deviation bounds explain why a large change can be unsafe, but they do not say when a finite rollout stops supporting an accurate update for the current policy. We study this transition through two ordered questions. First, when is the ordinary full-trajectory importance-sampling (IS) estimator still reliable? We show that normalized effective sample size (ESS) controls the policy-mismatch component of its finite-sample error and, through smoothness, yields a high-probability one-step improvement guarantee. Second, when does coefficient control provide a better update? Lower gradient-estimation MSE is not enough: the controlled estimator must also retain sufficient alignment with the population gradient. We capture this requirement with a full fixed-step improvement certificate. A Gaussian bandit instantiates the finite-moment ESS reliability bound, estimator crossover, and one-step policy improvement. Although TIS-3 has the lowest MSE along the displayed path, IS gives the better update at high ESS; only after effective support contracts does TIS-3 overtake it, while the more aggressive PPO mask never reaches the oracle envelope. In the reported large-model trajectories, ESS-conditioned safeguards preserve early learning while preventing late failure.

1 Introduction

Reinforcement learning has become central to post-training large language models [1, 15, 18]. In scalable training systems, policy optimization is often off-policy, a rollout policy generates responses that a trainer reuses for one or more gradient updates, so the current policy need not match the policy that produced the data [11].

Policy optimization in traditional applications require the control of the policy optimization, preventing from deviating too much from the existing methods. Proximal policy optimization has been central effective to limit the resulting policy deviation, PPO-style objectives, including GRPO, use ratio-based clipping or masking [13–15]. The theoretical understanding is based on the trust-region guarantees. Based on the population-level view, the too much policy deviation is not good.

Despite the widespread use of the theoretical understanding of trust region, the practical observation in the application of the large language models often show no-collapse or even better performance in less clipping, even the permissive updates like TIS or no-clips have shown to work. However, these estimators can learn rapidly, yet their success need not persist. Figure 1 shows TIS caps 3 and 5 improving for many updates before abruptly collapsing. This provokes the central theoretical puzzle: why can the same estimator be productive and unreliable at the same time?

In this paper, we shift the While classical trust-region theory predicts the policy to degrade rapidly,

Classical trust-region theory explains why moving too far from the rollout policy can be unsafe [11, 13, 14]. It does not, however, resolve the delayed transition in Figure 1: failure follows many




Figure 1: Delayed failure under permissive TIS updates on Qwen3-30B-A3B. Both TIS 3 and TIS 5 improve AIME-2024 mean@16 for many updates, an ESS collapse emerges, and performance subsequently fails. This ordering motivates the question of when a finite rollout stops supporting a reliable policy-gradient update.

apparently successful steps rather than one conspicuous move. Repeated updates can progressively concentrate complete-trajectory likelihood weights, reducing the effective support provided by the rollout even when the empirical update uses factorwise coefficients. A policy-deviation bound controls a population approximation; it does not track the statistical adequacy of the available rollout. The practical question is therefore: when does a finite rollout cease to support a reliable next update?

We answer this question in two stages. We first analyze the ordinary full-trajectory importance-sampling (IS) estimator, for which normalized effective sample size (ESS) replaces the nominal batch size N with the effective count $N\rho$ in the finite-sample error bound. Combined with smoothness, this yields an ESS-dependent one-step improvement guarantee. We then compare the factorwise coefficient-controlled estimators used in practice. Low ESS indicates when the IS guarantee weakens, but it does not determine which alternative update is better. Clipping or truncation can reduce sampling variation and MSE while also removing useful population-gradient alignment. We therefore derive a full fixed-step certificate that determines when the risk reduction is sufficient to justify the signal removed by control.

An analytically tractable Gaussian bandit then follows the theory in order. It first checks the finite-moment ESS reliability bound and safe-step guarantee, then computes the pairwise certificate crossovers of TIS and PPO against IS. The result isolates the missing term in an MSE-only comparison: TIS-3 has the lowest MSE throughout the displayed path, yet IS remains the better update at high ESS because it retains more useful alignment. Large-model trajectories establish the operational consequence. Sample ESS exposes support loss under both TIS 3 and an unclipped update, while ESS-conditioned safeguards preserve permissive learning and activate control only in the low-support regime. On Qwen3-30B-A3B, ungated TIS 3 collapses to 0.417%, whereas the ESS-conditioned TIS-3/Clip-Higher trajectory retains 44.8% at step 200.

1.1 Contributions

1. We show that normalized ESS controls the policy-mismatch component of the finite-sample error of the full-trajectory IS gradient estimator and yields an ESS-dependent one-step population-improvement guarantee.

2. We derive a fixed-step certificate that separates MSE reduction from lost population-gradient alignment and identifies when coefficient control provides the stronger update.
3. We instantiate the finite-moment ESS reliability bound in a Gaussian bandit, followed by exact certificate comparisons among IS, TIS, and PPO and one-step policy improvement; we then show in large-model training that ESS supplies an actionable timing signal for safeguards that prevent late failure.

2 Related work

Policy-deviation guarantees. Trust-region and proximal methods control policy movement to obtain stable improvement guarantees or practical update rules [11, 13, 14]. These analyses provide the standard explanation for why large policy changes can be unsafe. Our framework is complementary: it studies when a fixed finite rollout ceases to provide a reliable estimate of the current policy gradient. The distinction is between a population-level approximation guarantee and a finite-sample reliability question.

Importance sampling and effective support. Importance-sampling policy optimization uses likelihood-ratio moments and Rényi divergence to quantify the reliability of off-policy estimators [5, 6]. In these results, effective rather than nominal sample size controls estimation error. We apply the same principle directly to the vector-valued policy-gradient estimator and use normalized ESS as the regime coordinate.

Reliable policy-gradient updates. Smoothness-based analyses translate the error of a finite gradient estimate into a lower bound on policy improvement and use this relation to choose step sizes or sample sizes [8, 10]. We combine this optimization argument with an ESS-dependent off-policy reliability bound.

Stability in RLVR. Recent methods alter ratio boundaries, sequence weights, divergence controls, or update rules to improve the stability of large-model reinforcement learning [2, 3, 7, 9, 16, 18, 19]. Rather than proposing another static boundary, we study the statistical transition that makes a permissive update unreliable and then ask when an intervention can recover reliability.

3 Preliminaries

This section provides the background and notation used in the following analysis.

3.1 Policy-gradient estimators from a fixed rollout

Problem formulation. Let $X \sim \nu$ denote a context, $Q(\cdot | X)$ the rollout policy, and $P_\theta(\cdot | X)$ the current policy. A complete sample is $Z = (X, Y)$, where $Y = (Y_1, \dots, Y_T)$ and $R(Z)$ is its reward. For each factor t , define the likelihood ratio and score

$$r_{\theta,t}(Z) := \frac{P_\theta(Y_t | X, Y_{<t})}{Q(Y_t | X, Y_{<t})}, \quad \ell_{\theta,t}(Z) := \nabla_\theta \log P_\theta(Y_t | X, Y_{<t}). \quad (1)$$

Here t indexes a policy factor, such as a token or decision step. Because both policies factorize along the trajectory, the complete-sample likelihood ratio is

$$W_\theta(Z) := \frac{P_\theta(Y | X)}{Q(Y | X)} = \prod_{t=1}^T r_{\theta,t}(Z). \quad (2)$$

The learning objective is to maximize the expected reward $J(\theta) := \mathbb{E}_{P_\theta}[R]$. Let $A(Z) := R(Z) - b(X)$ be a baseline-centered reward, where $b(X)$ does not depend on Y . Under $P_\theta(\cdot | X) \ll Q(\cdot | X)$, change of measure expresses the population gradient under the rollout distribution:

$$g(\theta) := \nabla J(\theta) = \mathbb{E}_Q[W_\theta f_\theta], \quad f_\theta(Z) := A(Z) \sum_{t=1}^T \ell_{\theta,t}(Z) = A(Z) \nabla_\theta \log P_\theta(Y | X). \quad (3)$$

While trust-region arguments motivate safeguards such as TRPO and PPO-style clipping [13, 14], these arguments are typically made at the population level. Practical algorithms, however, do not observe $g(\theta)$ and must rely on its finite-sample counterpart. Given i.i.d. rollout samples $Z_i \sim \nu(dx)Q(dy | x)$, $i = 1, \dots, N$, the ordinary full-trajectory importance-sampling (IS) policy-gradient estimator is

$$\hat{g}_N(\theta) := \frac{1}{N} \sum_{i=1}^N W_{\theta,i} f_{\theta,i}. \quad (4)$$

The importance-weighted nature of this estimator can produce poor agreement between the finite-sample update and the population gradient [5, 6]. This exposes an important but often overlooked question: when is the finite-sample gradient reliable?

Choice of gradient estimator. To formalize this comparison, write $\mathbf{r}_\theta = (r_{\theta,1}, \dots, r_{\theta,T})$ and let $u = (u_1, \dots, u_T)$ be a collection of coefficient rules, with $u_t : \mathbb{R}_+^T \times \mathbb{R} \rightarrow \mathbb{R}$. Define

$$\psi_\theta(u; Z) := A(Z) \sum_{t=1}^T u_t(\mathbf{r}_\theta, A(Z)) \ell_{\theta,t}(Z), \quad \hat{g}_N(u; \theta) := \frac{1}{N} \sum_{i=1}^N \psi_\theta(u; Z_i). \quad (5)$$

Existing ratio-based policy-gradient estimators can be represented by particular choices of u . The IS estimator sets $u_{\text{IS},t} = W_\theta$ for every factor. One popular alternative is PPO/GRPO ratio clipping, whose clipped surrogate induces the following gradient-masking rule with radius $0 < \epsilon < 1$ [14, 15]:

$$u_{\text{P},t}(\mathbf{r}, A) := r_t [\mathbf{1}\{A \geq 0, r_t \leq 1 + \epsilon\} + \mathbf{1}\{A < 0, r_t \geq 1 - \epsilon\}]. \quad (6)$$

Such clipping rules create a conventional bias–variance tradeoff. More aggressive clipping can reduce sampling variation, but it can also increase bias by removing gradient terms. This raises the natural question: what governs the reliability of a gradient estimator, and when is clipping worth the signal it removes? We answer the first part by introducing Rényi divergence and effective sample size (ESS).

3.2 Gradient estimators and Rényi divergence

The dispersion of the importance weights indicates how dissimilar P_θ and Q are. Rényi divergence formalizes this notion as an information-theoretic dissimilarity between probability measures [12, 17], while importance-sampling analyses connect it to the finite-sample behavior of importance-weighted means [5, 6].

Rényi divergence is a family of measures indexed by a moment order $\alpha > 1$. For the current and rollout policies, define

$$D_\alpha(P_\theta \| Q) := \frac{1}{\alpha - 1} \log \mathbb{E}_Q[W_\theta^\alpha], \quad d_\alpha(P_\theta \| Q) := \exp\{D_\alpha(P_\theta \| Q)\}. \quad (7)$$

Larger values of α place progressively more emphasis on rare, large likelihood ratios. The order-two case is central because $\mathbb{E}_Q[W_\theta^2]$ controls standard MSE bounds for bounded importance-weighted contributions. Importantly, its reciprocal is the normalized effective sample size [5, 6]:

$$\rho_\theta := \frac{1}{d_2(P_\theta \| Q)} = \frac{1}{\mathbb{E}_Q[W_\theta^2]} = \exp\{-D_2(P_\theta \| Q)\}. \quad (8)$$

Its observable normalized sample analogue is

$$\hat{\rho}_\theta := \frac{\left(\sum_{i=1}^N W_{\theta,i}\right)^2}{N \sum_{i=1}^N W_{\theta,i}^2}. \quad (9)$$

It lies in $[1/N, 1]$ and is computed directly from the complete-trajectory weights in the batch used for the update. When $\mathbb{E}_Q[W_\theta^2] < \infty$, Jensen’s inequality and $\mathbb{E}_Q[W_\theta] = 1$ give $\rho_\theta \in (0, 1]$, and $N\rho_\theta$ is the effective sample count corresponding to nominal batch size N . As the second moment of the likelihood ratio grows, both ρ_θ and this effective count contract. The identity $\mathbb{E}_Q[(W_\theta - 1)^2] = \rho_\theta^{-1} - 1$ shows that ρ_θ is exactly the inverse second-moment measure of ratio dispersion.

Importance-sampling theory makes this quantity operational. If $\|f_\theta\|_\infty := \sup_z \|f_\theta(z)\|_2 < \infty$, the MSE of the IS estimator is at most $\|f_\theta\|_\infty^2 / (N\rho_\theta)$. Normalized ESS therefore captures the policy-mismatch component of IS reliability. The next section determines when that reliability is sufficient for improvement and then asks whether coefficient control offers a better update.

4 Theoretical Analysis

This section answers the preceding questions in order. We first connect gradient-estimator reliability to policy improvement, then use normalized ESS to control the reliability of the IS estimator, and finally ask when clipping provides the stronger update.

4.1 Gradient-estimator reliability and policy improvement

The preceding section distinguishes the finite-sample IS estimate from the population gradient it targets. Here, we connect estimator reliability to a lower bound on policy improvement. The following lemma formalizes this connection for an arbitrary gradient estimator.

Lemma 1 (A reliable gradient estimate yields policy improvement). *Let $\hat{g}$ be an estimate of $g = \nabla J(\theta)$ satisfying $\|\hat{g} - g\|_2 \leq \varepsilon$. Suppose J is L -smooth along the update segment for $L > 0$, and define $\gamma = L^{-1}(1 - \varepsilon/\|\hat{g}\|_2)_+$, with $\gamma = 0$ when $\hat{g} = 0$. Then*

$$J(\theta + \gamma\hat{g}) - J(\theta) \geq \frac{1}{2L} (\|\hat{g}\|_2 - \varepsilon)_+^2. \quad (10)$$

This deterministic lemma uses the same smoothness mechanism as existing policy-gradient analyses [8, 10]. Its implication is direct: the error radius measures how far the finite-sample direction can lie from the population gradient. When this radius is smaller than $\|\hat{g}\|_2$, the estimator yields a positive lower bound on policy improvement.

ESS and the IS update. However, the population gradient is unknown, so this reliability cannot be evaluated directly. Normalized ESS provides the bridge: for the IS estimator, it converts likelihood-ratio dispersion into a bound on gradient-estimation error. At update k , condition on the preceding updates and write $W_k := W_{\theta_k}$, $f_k := f_{\theta_k}$, $g_k := g(\theta_k)$, $\hat{g}_k := \hat{g}_N(\theta_k)$, and $\rho_k := \rho_{\theta_k}$. Assume that the N rollout samples are independent draws from Q , $d_2(P_{\theta_k} \| Q) < \infty$, and $\mathbb{E}_Q[W_k^2 \| f_k \|_2^2] < \infty$. Define the importance-weighted contribution scale

$$M_{2,k} := \frac{\mathbb{E}_Q[W_k^2 \| f_k \|_2^2]}{\mathbb{E}_Q[W_k^2]} = \rho_k \mathbb{E}_Q[W_k^2 \| f_k \|_2^2]. \quad (11)$$

This factorizes the weighted second moment as $M_{2,k}/\rho_k$: ρ_k records ratio dispersion, while $M_{2,k}$ also captures its coupling with the score-advantage magnitude.

Because the IS estimator is unbiased, its exact mean-squared error is

$$\mathbb{E} \|\hat{g}_k - g_k\|_2^2 = \frac{1}{N} \left\{ \frac{M_{2,k}}{\rho_k} - \|g_k\|_2^2 \right\}. \quad (12)$$

Consequently,

$$\mathbb{E} \|\hat{g}_k - g_k\|_2^2 \leq \frac{M_{2,k}}{N\rho_k}. \quad (13)$$

Likelihood-ratio dispersion therefore enters through the effective count $N\rho_k$, replacing the nominal sample size N , while $M_{2,k}$ records the remaining contribution scale. If f_k is bounded, then $M_{2,k} \leq \|f_k\|_\infty^2$, recovering the familiar bounded-integrand form. Equation (13) controls the mean-squared error; the next lemma converts this result into a high-probability error radius for the realized update.

Lemma 2 (Normalized ESS controls the gradient-error radius). *Under the preceding setup, for any $0 < \delta \leq 1$, with probability at least $1 - \delta$,*

$$\|\hat{g}_k - g_k\|_2 \leq \sqrt{\frac{M_{2,k}}{\delta N \rho_k}}. \quad (14)$$

Holding the contribution scale fixed, the radius grows as $(N\rho_k)^{-1/2}$. Thus, low ESS weakens the reliability guarantee for the IS gradient, consistent with the motivating example in the introduction. We now combine this radius with Lemma 1 to obtain an ESS-dependent improvement guarantee.

Theorem 3 (ESS-dependent improvement for the IS estimator). *Under the conditions of Lemma 2, suppose that J is L_k -smooth along the update segment for $L_k > 0$. For any $0 < \delta \leq 1$, define*

$$\varepsilon_k(\delta) := \sqrt{\frac{M_{2,k}}{\delta N \rho_k}}, \quad \gamma_k := \frac{1}{L_k} \left(1 - \frac{\varepsilon_k(\delta)}{\|\hat{g}_k\|_2} \right)_+, \quad (15)$$

with $\gamma_k = 0$ when $\hat{g}_k = 0$. Then, with probability at least $1 - \delta$,

$$J(\theta_k + \gamma_k \hat{g}_k) - J(\theta_k) \geq \frac{1}{2L_k} (\|\hat{g}_k\|_2 - \varepsilon_k(\delta))_+^2. \quad (16)$$

Theorem 3 makes the role of ESS explicit. Its lower bound is positive when

$$\|\hat{g}_k\|_2 > \sqrt{\frac{M_{2,k}}{\delta N \rho_k}}.$$

Thus, the IS update remains certifiably useful while effective support is sufficiently large relative to the observed gradient signal. Holding the other terms fixed, both the admissible step and the certified improvement shrink as ESS falls. This complements population policy-deviation guarantees, which characterize the consequences of changing the policy, by controlling the finite-rollout direction used to make that change. Having established when the IS update is reliable, we next ask when an alternative gradient estimator provides the stronger one-step certificate.

4.2 Choosing an alternative gradient estimator

The preceding result connects the reliability of the IS estimator to policy improvement. When this guarantee weakens, one may instead use a controlled coefficient rule, such as truncated IS or PPO-style masking. We compare a general rule u with IS in two steps: first as a question of gradient estimation, and then as a question of policy improvement at a common fixed step. Write $\hat{g}_k(u) := \hat{g}_N(u; \theta_k)$. For a rule u , adding and subtracting the estimator mean gives the MSE decomposition

$$\begin{aligned} m_k(u) &:= \mathbb{E} \|\hat{g}_k(u) - g_k\|_2^2 \\ &= \|\mathbb{E}[\hat{g}_k(u)] - g_k\|_2^2 + \mathbb{E} \|\hat{g}_k(u) - \mathbb{E}[\hat{g}_k(u)]\|_2^2. \end{aligned} \quad (17)$$

The two terms are squared bias and sampling variation. Since the IS estimator is unbiased, a controlled rule has lower MSE exactly when its reduction in sampling variation exceeds the squared bias introduced by changing the coefficients. We call this the estimator-level bias–variance crossover. The next theorem gives the expected-improvement certificate used to compare such estimators at a common fixed step.

Theorem 4 (Improvement certificate). *Suppose that J is L_k -smooth on every update segment generated by a coefficient rule u almost surely, for a constant $L_k > 0$, and that $\mathbb{E} \|\hat{g}_k(u)\|_2^2 < \infty$. For any step size $\eta_k > 0$,*

$$\mathbb{E}[J(\theta_k + \eta_k \hat{g}_k(u)) - J(\theta_k)] \geq \eta_k g_k^\top \mathbb{E}[\hat{g}_k(u)] - \frac{L_k \eta_k^2}{2} \mathbb{E} \|\hat{g}_k(u)\|_2^2. \quad (18)$$

Denote the right-hand side by

$$B_k(u; \eta_k) := \eta_k g_k^\top \mathbb{E}[\hat{g}_k(u)] - \frac{L_k \eta_k^2}{2} \mathbb{E} \|\hat{g}_k(u)\|_2^2. \quad (19)$$

The certificate has two parts. The first rewards alignment with the population gradient; the second penalizes the expected squared update under smoothness. Coefficient control improves the certificate only when the penalty it removes outweighs the useful alignment it removes. The next corollary states this crossover directly.

Corollary 5 (IS-to-controlled-estimator crossover). *Let u_{IS} denote the IS coefficient rule and let u be any alternative rule. Suppose the same deterministic L_k is valid almost surely for both families of update segments. With a common step size $0 < \eta_k \leq 1/L_k$, the alternative rule has the stronger certificate if and only if*

$$B_k(u; \eta_k) > B_k(u_{\text{IS}}; \eta_k) \iff m_k(u_{\text{IS}}) - m_k(u) > (1 - L_k \eta_k) \{ \mathbb{E} \|\hat{g}_k(u_{\text{IS}})\|_2^2 - \mathbb{E} \|\hat{g}_k(u)\|_2^2 \}. \quad (20)$$

When the alternative lowers the update second moment, the right-hand side of Equation (20) is its discounted reduction. Its MSE gain must exceed this amount. Since the IS estimator is unbiased, the same condition has the equivalent alignment form

$$\|g_k\|_2^2 - g_k^\top \mathbb{E}[\hat{g}_k(u)] < \frac{L_k \eta_k}{2} \{ \mathbb{E} \|\hat{g}_k(u_{\text{IS}})\|_2^2 - \mathbb{E} \|\hat{g}_k(u)\|_2^2 \}. \quad (21)$$

The left-hand side is the population-gradient alignment lost by changing the estimator; the right-hand side is the reduction in the quadratic smoothness penalty.

Remark 6 (Truncated importance sampling). Upper truncation, $u_{T,t} = \min\{r_t, c\}$, caps each selected ratio without deleting its tail contribution. It is therefore a cheap variance safeguard with comparatively mild bias intervention. Its control is also limited: it acts only on the chosen upper tail, so the reduction in the update second moment can be modest under sequence-level mismatch.

Remark 7 (PPO-style masking). The PPO rule in Equation (6) removes saturated positive-advantage terms and low-ratio negative-advantage terms. It can reduce the update second moment more aggressively than upper truncation, but it also discards more population-gradient alignment. Equations (20) and (21) state when that stronger intervention is worthwhile.

5 Policy optimization and ESS

This section makes the theoretical distinction observable in two settings. An analytically tractable Gaussian bandit first instantiates IS reliability, estimator crossover, and one-step policy improvement. Large-model trajectories then show how observed effective support evolves around delayed optimization failure.

5.1 Policy optimization in a Gaussian bandit

The Gaussian bandit makes our two-stage argument fully observable. We hold the current policy fixed and move the rollout policy farther away, which creates a monotone path from high to low normalized ESS. A quadratic reward then makes the population gradient, estimator moments, and expected one-step improvement available in closed form. We can therefore test the complete chain: ESS first predicts when IS loses finite-sample reliability, and the gain comparison then selects the estimator that makes the best use of the remaining support. Appendix A.1 gives the calculations and protocol.




Figure 2: **Stage I: ESS predicts IS reliability and certified improvement.** We use $g = 2$, $N = 32$, and $1 - \delta = 0.9$. (a) Exact IS gradient MSE and the finite-moment upper bound in Equation (13). (b) Positive Theorem 3 certificates and actual positive improvement under the corresponding adaptive step.

The first stage asks when an IS update remains reliable. Figure 2(a) answers this question at the estimator level: the exact IS MSE stays below the bound from Lemma 2, and both rise sharply as ESS falls. Theorem 3 uses this same error radius to build an adaptive step, so estimator reliability should translate directly into policy improvement. Panel (b) confirms that translation. Across the full ESS path, the probability of actual improvement tracks the probability of a positive certificate with a maximum gap of only 0.00891. As support contracts, both probabilities fall together. ESS therefore identifies when a reliable IS update is available, but this first stage leaves one question open: which controlled estimator should be used after IS loses its advantage?




Figure 3: **Stage II: the gain comparison selects the estimator.** We compare IS, TIS with $c = 3$, and PPO masking with $\epsilon = 0.2$ at $\eta = 0.4$. (a) Exact expected gain relative to IS; positive values favor the controlled estimator. (b) Exact expected gain around the two pairwise crossovers. Vertical lines and open markers identify the intersections.

The second stage resolves that choice through the estimator-selection criterion in Corollary 5. Figure 3(a) plots the resulting expected-gain difference relative to IS. TIS-3 has the lowest MSE throughout the path, yet IS retains more useful gradient alignment and therefore has the larger gain at high ESS. As ESS falls, the variance reduction from truncation becomes decisive, and TIS-3 overtakes IS at $\rho = 0.105$. PPO removes still more alignment, so it crosses IS only at $\rho = 0.0392$. Panel (b) makes both transitions visible as intersections of the exact policy-improvement curves. The resulting upper envelope selects IS in the high-support regime and TIS-3 in the low-support regime, exactly as predicted by the gain criterion.

Together, Figures 2 and 3 turn our theory into an operational principle. ESS determines when coefficient control becomes valuable, while the gain criterion determines which control preserves the strongest update. The remaining question is whether ESS provides this timing signal during large-model training, where the policy pair evolves after every update.

5.2 Policy optimization in large language models

Large-model training turns the fixed-policy comparison into a temporal one: ESS must fall early enough to trigger control before the permissive trajectory fails. We test this requirement on the ungated TIS-3 trajectory from Figure 1 and an unclipped Qwen3-30B-A3B trajectory. Both are evaluated on AIME-2024 using mean@16. Figure 4 aligns raw sequence ESS, its trailing seven-update average, and held-out performance along each training path.




Figure 4: **Sequence ESS provides an early training-time signal.** Thin blue curves show raw ESS, thick blue curves show its trailing seven-update average, and green curves show AIME-2024 mean@16. Blue dotted lines mark the first sustained sub-0.1 ESS episode; green dashed lines mark peak validation. The gray TIS-3 segment occurs after collapse.

Figure 4 shows the same temporal ordering in both permissive updates. For the unclipped trajectory, the first sustained sub-0.1 ESS episode begins near step 103, 37 updates before validation reaches its peak. Under TIS 3, the corresponding episode begins near step 71, while validation remains close to its maximum and before the subsequent collapse. The later rebound in TIS-3 ESS occurs only after performance has vanished, which distinguishes renewed policy overlap from recovered policy quality. In both cases, sustained support contraction is visible before the failure that motivates intervention.

These trajectories provide the training-time counterpart of the Gaussian reliability result. Sequence ESS supplies an observable trigger for the low-support regime, and the gain comparison supplies the rule used after that trigger. This combination yields our policy-optimization design: retain the permissive estimator while ESS is high, then activate a prespecified safeguard when ESS becomes low. Section 6 tests whether this timing rule preserves early learning while preventing late collapse.

6 ESS-conditioned policy optimization

Section 5 identifies ESS as an observable coordinate for effective support. We now ask whether acting on this signal resolves the central optimization tradeoff: retain a permissive update while the rollout remains informative, then activate control as effective support contracts. We first test this rule in Qwen3-30B-A3B math RLVR and then change the model scale and training objective.

6.1 Update rule and protocol

Let $\hat{g}_{\text{perm}}$ denote the permissive update and $\hat{g}_{\text{safe}}$ a prespecified safeguard. Given a threshold τ , the ESS-conditioned update is

$$\hat{g}_{\text{gate}} = \begin{cases} \hat{g}_{\text{perm}}, & \hat{\rho} \geq \tau, \\ \hat{g}_{\text{safe}}, & \hat{\rho} < \tau. \end{cases} \quad (22)$$

ESS determines when the update changes; the selected safeguard determines how the low-support update is controlled. We fix $\tau = 0.1$ for all primary comparisons and vary it only in the threshold

ablation. This value was prescribed for the language-model runs and is not derived from the Gaussian thresholds in Section 5.1. GRPO, DPPO, and the unclipped construction use the raw $\hat{\rho}$ in Equation (9). Clip-Higher instead uses the analogous shaped ESS computed from post-cap sequence coefficients.

Evaluation protocol. The main experiment trains Qwen3-30B-A3B for up to 200 policy updates on the boxed mathematical-reasoning data of Yu et al. [18] and evaluates AIME-2024 using mean@16. The primary permissive branch is TIS with ratio cap 3; an explicit ablation instead uses the unclipped update. At low ESS, we test three prespecified safeguards: conventional GRPO clipping [15], Clip-Higher [18], and DPPO [11]. The static 30B recipes use their native token-mean aggregation, whereas the TIS-3 conditional recipes use sum normalization. Panels (a)–(c) of the main figure therefore compare complete recipes; the TIS-3 pair in panel (d) and the unclipped comparison isolate ESS conditioning.

6.2 Main results

Figure 5 establishes the central empirical result. Each conditional recipe uses TIS 3 while effective support is high and activates the named safeguard only when normalized ESS falls below 0.1. Across GRPO, Clip-Higher, and DPPO, the conditional recipes achieve both a higher peak and a higher final AIME score than the corresponding static recipes, with endpoint gains ranging from 6.88 to 12.3 percentage points. The matched TIS-3 comparison isolates the failure that conditioning prevents: ungated TIS 3 peaks at 31.7% and then collapses to 0.417%, whereas the ESS-conditioned trajectory reaches 47.3% and retains 44.8% at step 200. ESS conditioning therefore does more than stabilize a permissive estimator. It preserves its learning advantage while support is broad and introduces control only after effective support contracts.




Figure 5: ESS conditioning combines permissive high-support learning with low-support control on Qwen3-30B-A3B. Each conditional trajectory uses TIS 3 when $\hat{\rho} \geq 0.1$ and the named safeguard otherwise. Panels (a)–(c) compare the complete conditional and static recipes; the former use sum normalization and the latter their native token-mean aggregation. Panel (d) gives the matched collapse-prevention comparison against ungated TIS 3, using the TIS-3/Clip-Higher trajectory from panel (b). Direct labels report the final observed AIME-2024 mean@16.

Figure 6 uses the unclipped trajectory diagnosed in Figure 4 as the permissive reference. Conditional clipping preserves a comparable peak and raises step-200 performance from 34.6% to

38.1%. Thus, the same ESS signal can retain most of the permissive estimator’s learning while limiting late-stage degradation without relying on a particular truncation cap.




Figure 6: The timing principle extends to a fully unclipped high-support update. Conventional clipping is activated only when raw normalized ESS is below 0.1. Conditional control preserves a comparable peak while reducing the late performance loss. Direct labels report step-200 AIME-2024 mean@16.

Results on Dense models Figure 7 changes from the 30B mixture-of-experts model to dense Qwen3-8B while keeping math RLVR fixed. The ESS-conditioned construction finishes above its static counterpart in both aggregation-matched comparisons. The gains are smaller than the 30B rescue, but the endpoint ordering persists under both Clip-Higher and DPPO. This extends the support-conditioned design beyond the model on which the failure was first observed.




Figure 7: Dense-model transfer on Qwen3-8B. Panel (a) compares token-mean Clip-Higher with its token-mean ESS-conditioned construction. Panel (b) compares static and ESS-conditioned DPPO under the same sum-normalized aggregation. Direct labels give the step-200 AIME-2024 mean@16.

Results on RLHF We next change both the objective and the data, training Qwen3-4B-Instruct-2507 for 200 updates on Anthropic HH-RLHF [1] using Skywork-Reward-Llama-3.1-8B [4] as the reward model. Figure 8 shows that the complete TIS-3-plus-ESS recipe reaches a reward-model score

of 60.7, compared with 43.0 for GRPO. The advantage therefore survives a shift from verifiable math rewards to learned-reward optimization. This experiment measures optimization of the training reward model rather than human preference directly.




Figure 8: Transfer to open-ended RLHF on Qwen3-4B-Instruct-2507. Thin curves are per-step reward-model scores; thick curves are trailing seven-step means. Direct labels give the step-200 observations. This is a complete-recipe comparison between GRPO and TIS 3 with an ESS-conditioned safeguard.

Across these experiments, ESS contributes timing rather than a new clipping rule: it converts conservative control from a permanent constraint into a targeted intervention.

6.3 Ablations and alternative interventions

Collapse prevention does not depend on one exact threshold. Figure 9 varies only the ESS threshold for the TIS-3/GRPO construction. Thresholds 0.05, 0.1, and 0.2 all avoid the catastrophic 0.417% endpoint of ungated TIS 3, with last-reported scores of 43.1%, 38.5%, and 34.8%, respectively. Performance changes because the threshold determines how long the permissive branch remains active. Thus, τ is an operating point for the ESS diagnostic rather than a universal theoretical constant; the failure-prevention effect persists across the tested fourfold range.

Alternative low-support actions separate detection from control. ESS can also trigger a complete update veto rather than clipping. Figure 10 shows that this simpler action rescues TIS 5 from a 0.00% collapse to a 39.4% endpoint. Skipping is more conservative when the permissive trajectory remains viable: with the unclipped update it finishes at 35.6%, below the 38.1% obtained by ESS-conditioned clipping, and sacrifices more of the attainable peak. Across both experiments, ESS provides a common trigger for the low-support regime, but the response still matters. The stronger performance under conditional clipping is consistent with controlled estimation retaining useful signal that a full update veto discards.




Figure 9: Sensitivity to the ESS threshold in the Qwen3-30B-A3B TIS-3/GRPO construction. All three thresholds prevent the ungated TIS-3 collapse; changing τ shifts the balance between the permissive and safeguarded branches. Labels report peak and last-observed AIME-2024 mean@16.




Figure 10: ESS-conditioned skipping as an alternative low-support action. Panel (a) shows that a hard veto stabilizes an unclipped update but sacrifices more learning than ESS-conditioned clipping. Panel (b) gives the matched TIS-5 comparison, in which skipping prevents the static trajectory’s collapse. All curves report AIME-2024 mean@16.

Together, the threshold and action ablations show that collapse prevention comes from conditioning control on effective support, while retained learning depends on how aggressively the low-support update is modified.

7 Conclusion

We studied clipping as a finite-sample estimator choice rather than a default property of policy optimization. The theory separates gradient-estimation error from the useful update retained after control: reducing MSE improves the fixed-step certificate only when that gain compensates for the signal removed by clipping. This distinction makes IS preferable while effective support is broad and explains why controlled estimators become attractive as mismatch grows. The Gaussian construction instantiates these comparisons and shows why mild truncation can dominate more aggressive masking. The reported language-model trajectories then show that sequence ESS can reveal support contraction and time a prespecified safeguard without serving as the oracle itself. In the RLVR and RLHF examples, the resulting conditional constructions retain permissive learning in the higher-support regime and protect training more strongly as support contracts. ESS therefore provides the operational link between importance-weight reliability and the decision of when to control an update.

A Numerical Experiments

A.1 Gaussian bandit

This section gives the calculations used in Section 5.1.

Setup We compare IS, TIS with cap $c = 3$, and the PPO-style rule in Equation (6), using the exact current-policy value baseline.

Because the score contribution is polynomial and each mask boundary is linear in z , the means, MSEs, and update second moments are finite sums of truncated Gaussian moments.

We choose a quadratic reward because it makes the smoothness certificate equal to the true expected gain. Any separation between the two oracles is therefore genuine rather than an artifact of a loose bound.

Consider the current policy and rollout policy

$$P_\theta = \mathcal{N}(\theta, 1), \quad Q_\Delta = \mathcal{N}(\theta - \Delta, 1), \quad \Delta \geq 0, \quad (23)$$

with reward $R(a) = -\frac{1}{2}(a - a^*)^2$. We hold the current policy fixed and vary only the rollout gap Δ . Writing $g := a^* - \theta > 0$, the objective is

$$J(\theta) = -\frac{1}{2}\{(\theta - a^*)^2 + 1\}, \quad \nabla J(\theta) = g, \quad L = 1. \quad (24)$$

Here θ is the fixed pre-update mean and Q_Δ remains frozen during the hypothetical one-step update. We vary Δ only across common-state comparisons. For $z = a - \theta$, the likelihood ratio and normalized ESS reduce to

$$W_\Delta(z) = \exp\{\Delta z + \Delta^2/2\}, \quad \rho(\Delta) = \{\mathbb{E}_{Q_\Delta}[W_\Delta^2]\}^{-1} = \exp(-\Delta^2). \quad (25)$$

Normalized ESS is therefore a one-to-one coordinate for this branch.

Write $\mu_u := \mathbb{E}[\hat{g}_u]$, $s_u := \mathbb{E}[\hat{g}_u^2]$, and $B_u(\eta) := \eta g \mu_u - \eta^2 s_u / 2$. For this quadratic objective, the fixed-step certificate is exact:

$$\mathbb{E}[J(\theta + \eta \hat{g}_u) - J(\theta)] = \eta g \mu_u - \frac{\eta^2}{2} s_u = B_u(\eta). \quad (26)$$

We use $N = 32$, the standard PPO radius $\epsilon = 0.2$, TIS cap $c = 3$, and choose $g = 2$ and $\eta = 0.4$ so that the reliability and estimator-crossover regimes are visible on one monotone ESS path.

Exact Closed Form Computations Let $z = a - \theta$. Under the current policy, $z \sim \mathcal{N}(0, 1)$ and the score is z . The exact current-policy value baseline gives

$$A_g(z) = gz + \frac{1 - z^2}{2}, \quad f_g(z) := z A_g(z) = gz^2 + \frac{z - z^3}{2}. \quad (27)$$

Although f_g is unbounded, all polynomial Gaussian moments used below are finite, so the finite-moment conditions of Lemma 2 and Theorem 3 hold. IS uses $W_\Delta f_g$ and TIS uses $\min\{W_\Delta, c\} f_g$. The PPO-style candidate uses $W_\Delta f_g M_{\epsilon, \Delta}$, where

$$M_{\epsilon, \Delta}(z) := \mathbf{1}\{A_g(z) \geq 0, W_\Delta(z) \leq 1 + \epsilon\} + \mathbf{1}\{A_g(z) < 0, W_\Delta(z) \geq 1 - \epsilon\}. \quad (28)$$

Let $M_{\text{IS}} \equiv 1$ and $M_{\text{P}} = M_{\epsilon, \Delta}$. Change of measure gives the exact per-sample moments

$$\begin{aligned} \mu_u(\Delta) &= \mathbb{E}_{U \sim \mathcal{N}(0, 1)}[f_g(U) M_u(U)], \\ h_u(\Delta) &= e^{\Delta^2} \mathbb{E}_{V \sim \mathcal{N}(\Delta, 1)}[f_g(V)^2 M_u(V)], \quad u \in \{\text{IS}, \text{P}\}. \end{aligned} \quad (29)$$

For TIS, define $b_c(\Delta) := (\log c - \Delta^2/2)/\Delta$. Its exact moments are

$$\begin{aligned} \mu_{\text{T}}(\Delta) &= \mathbb{E}_{U \sim \mathcal{N}(0, 1)}[f_g(U) \mathbf{1}\{U \leq b_c\}] + c \mathbb{E}_{U \sim \mathcal{N}(-\Delta, 1)}[f_g(U) \mathbf{1}\{U > b_c\}], \\ h_{\text{T}}(\Delta) &= e^{\Delta^2} \mathbb{E}_{V \sim \mathcal{N}(\Delta, 1)}[f_g(V)^2 \mathbf{1}\{V \leq b_c\}] + c^2 \mathbb{E}_{V \sim \mathcal{N}(-\Delta, 1)}[f_g(V)^2 \mathbf{1}\{V > b_c\}]. \end{aligned} \quad (30)$$

In particular, $\mu_{\text{IS}} = g$. For an iid estimator with batch size N ,

$$m_u = (\mu_u - g)^2 + \frac{h_u - \mu_u^2}{N}, \quad s_u = \mu_u^2 + \frac{h_u - \mu_u^2}{N}. \quad (31)$$

For $g = 2$, the IS weighted-contribution scale in Equation (11) is

$$\begin{aligned} M_2(\Delta) &= \rho(\Delta) h_{\text{IS}}(\Delta) \\ &= \frac{\Delta^6}{4} - 2\Delta^5 + \frac{29\Delta^4}{4} - 18\Delta^3 + \frac{65\Delta^2}{2} - 24\Delta + \frac{29}{2}. \end{aligned} \quad (32)$$

Hence the exact IS MSE is $\{e^{\Delta^2} M_2(\Delta) - 4\}/32$ and its Lemma 2 envelope is $e^{\Delta^2} M_2(\Delta)/32$. Figure 2(b) uses 10^5 independent batches at each of 81 policy pairs and $\delta = 0.1$ to estimate the probability of a positive adaptive-step certificate and the probability of actual positive policy improvement. On the same draws, the minimum realized coverage of both the Lemma 2 event and the Theorem 3 inequality is 0.98277, above the required 0.9. Both $M_2(\Delta)$ and the adaptive step are evaluated from population quantities in this controlled theory validation; they are not proposed as an empirical gate.

The advantage changes sign at $g \pm \sqrt{g^2 + 1}$, and the PPO ratio boundaries are $(\log(1 \pm \epsilon) - \Delta^2/2)/\Delta$. Equations (29) and (30) are therefore finite sums of truncated normal polynomial moments. With $I_j(a, b) := \int_a^b z^j \phi(z) dz$, integration by parts gives

$$I_j(a, b) = a^{j-1} \phi(a) - b^{j-1} \phi(b) + (j-1) I_{j-2}(a, b), \quad j \geq 2, \quad (33)$$

starting from $I_0 = \Phi(b) - \Phi(a)$ and $I_1 = \phi(a) - \phi(b)$. Boundary terms vanish at infinity. Substituting these moments into Equation (31) and the exact gain in Equation (26) gives the reported boundaries.

Results We evaluate $\sqrt{-\log(0.9)} \leq \Delta \leq \sqrt{-\log(0.005)}$. For each $u \in \{\text{T}, \text{P}\}$, substitution into Corollary 5 gives

$$C_u(\rho) = \{m_{\text{IS}} - m_u\} - (1 - \eta)\{s_{\text{IS}} - s_u\} = \frac{2}{\eta}\{B_u - B_{\text{IS}}\}. \quad (34)$$

Numerical evaluation of the closed-form moments gives

$$\rho_{\text{IS}, \text{TIS3}} = 0.105, \quad \rho_{\text{IS}, \text{PPO}} = 0.0392. \quad (35)$$

At the TIS-3 root, the MSE reduction and discounted second-moment reduction both equal 3.94; at the PPO root, both equal 11.0. TIS-3 has lower MSE than both IS and PPO throughout the displayed branch, yet IS has the larger certificate for $\rho > 0.105$. Below that point TIS-3 forms the upper certificate envelope. PPO crosses IS at 0.0392 only pairwise: it remains below TIS-3 throughout the branch. Equation (26) makes the same ordering exact for expected one-step policy improvement.

B Additional coefficient-rule details

The main text treats coefficient control generically because different rules alter different parts of the gradient. A factorwise detached upper truncation uses $u_{\text{T},t}(\mathbf{r}, A) := \min\{r_t, c\}$. Its contribution obeys

$$\|\psi_k(u_{\text{T}}; Z)\|_2 \leq c|A(Z)| \sum_{t=1}^T \|\ell_{\theta_k,t}(Z)\|_2, \quad (36)$$

so it has a direct second-moment envelope whenever the right-hand side has a finite second moment.

PPO masking is different. For $A \geq 0$ and $c = 1 + \epsilon$,

$$0 \leq r - \min\{r, c\} = (r - c)_+ \leq r \mathbf{1}\{r > c\} = r - rM_\epsilon(r, A), \quad (37)$$

where $M_\epsilon(r, A) = \mathbf{1}\{A \geq 0, r \leq 1 + \epsilon\} + \mathbf{1}\{A < 0, r \geq 1 - \epsilon\}$. Truncation retains a bounded positive coefficient above the cap, whereas PPO removes that contribution. On the negative-advantage lower branch, PPO masks small ratios whereas upper truncation does not. A complete-sample mask is another distinct rule: it sets $u_t(\mathbf{r}, A) = W_k M_\epsilon(W_k, A)$ for every factor. We state which rule is used in each experiment rather than treating these implementations as interchangeable.

C Additional RLVR diagnostics

C.1 Intervention and training dynamics

The main text reports evaluation performance. We record here how often the conditional rule intervenes and the accompanying training dynamics.

Figure 11 shows the intended temporal ordering: intervention is already possible early, but its frequency increases as ESS contracts. The remaining figures record response length, entropy, and reward diagnostics.




Figure 11: Realized ESS-conditioned intervention schedules. Each panel pairs the per-step intervention fraction (thin black) and its trailing seven-step mean (thick black) with normalized ESS (red). GRPO and DPPO use raw ESS; Clip-Higher uses shaped ESS after the coefficient cap. The dotted line marks the threshold 0.1.




Figure 12: Training dynamics for the main comparisons. The static TIS 3 run develops extreme response length and entropy together with a collapse in training reward. The safeguards and their ESS-conditional variants avoid this failure over the reported 200 steps.




Figure 13: Detailed TIS 3 comparison. Under static TIS, ESS falls below 0.01 after step 77 and the response length and entropy diverge. The fixed-0.1 intervention increases the clipped fraction as ESS declines and keeps both diagnostics controlled.

D Proofs for the theoretical results

At update k , all expectations in this appendix are taken after conditioning on the preceding updates. The current policy is therefore fixed and the unused sampling units are independent draws from Q . We suppress this conditioning throughout.

D.1 Policy-gradient identity and mean-square bound

By definition, $f_k = (R - b(X))\nabla_\theta \log P_{\theta_k}(Y | X)$ and $W_k = dP_{\theta_k}/dQ$. Change of measure gives

$$\begin{aligned}\mathbb{E}_Q[W_k f_k] &= \mathbb{E}_{P_{\theta_k}}[(R - b(X))\nabla_\theta \log P_{\theta_k}(Y | X)] \\ &= \mathbb{E}_{P_{\theta_k}}[R\nabla_\theta \log P_{\theta_k}(Y | X)] \\ &\quad - \mathbb{E}_{P_{\theta_k}}[b(X)\nabla_\theta \log P_{\theta_k}(Y | X)].\end{aligned}\tag{38}$$

For the baseline term, condition further on X . Under the standard regularity condition that permits differentiation under the conditional integral,

$$\begin{aligned}\mathbb{E}_{P_{\theta_k}}[b(X)\nabla_\theta \log P_{\theta_k}(Y | X) | X] &= b(X) \int P_{\theta_k}(y | X) \nabla_\theta \log P_{\theta_k}(y | X) d\lambda(y) \\ &= b(X) \nabla_\theta \int P_{\theta_k}(y | X) d\lambda(y) \\ &= b(X) \nabla_\theta 1 = 0.\end{aligned}\tag{39}$$

The remaining term is the score-function identity. Hence

$$\mathbb{E}_Q[W_k f_k] = \nabla J(\theta_k) = g_k.\tag{40}$$

Averaging N independent copies gives $\mathbb{E}[\hat{g}_k] = g_k$.

For the mean-square calculation, define $\xi_{k,i} := W_{k,i} f_{k,i} - g_k$. Equation (40) implies $\mathbb{E}[\xi_{k,i}] = 0$. Since $\hat{g}_k - g_k = N^{-1} \sum_{i=1}^N \xi_{k,i}$,

$$\mathbb{E}\|\hat{g}_k - g_k\|_2^2 = \frac{1}{N^2} \sum_{i=1}^N \mathbb{E}\|\xi_{k,i}\|_2^2 + \frac{1}{N^2} \sum_{i \neq j} \mathbb{E}[\xi_{k,i}^\top \xi_{k,j}].\tag{41}$$

For $i \neq j$, independence and zero means give

$$\mathbb{E}[\xi_{k,i}^\top \xi_{k,j}] = \mathbb{E}[\xi_{k,i}]^\top \mathbb{E}[\xi_{k,j}] = 0.\tag{42}$$

The diagonal terms are identical, so

$$\mathbb{E}\|\hat{g}_k - g_k\|_2^2 = \frac{1}{N} \mathbb{E}_Q \|W_k f_k - g_k\|_2^2.\tag{43}$$

Expanding the remaining squared norm yields

$$\begin{aligned}\mathbb{E}_Q \|W_k f_k - g_k\|_2^2 &= \mathbb{E}_Q [W_k^2 \|f_k\|_2^2] - 2g_k^\top \mathbb{E}_Q [W_k f_k] + \|g_k\|_2^2 \\ &= \mathbb{E}_Q [W_k^2 \|f_k\|_2^2] - \|g_k\|_2^2,\end{aligned}\tag{44}$$

where the last equality uses Equation (40). Combining Equations (43) and (44) gives

$$\mathbb{E}\|\hat{g}_k - g_k\|_2^2 = \frac{1}{N} \{ \mathbb{E}_Q[W_k^2 \|f_k\|_2^2] - \|g_k\|_2^2 \}. \quad (45)$$

Finally, Equation (11) implies

$$\begin{aligned} \mathbb{E}\|\hat{g}_k - g_k\|_2^2 &\leq \frac{1}{N} \mathbb{E}_Q[W_k^2 \|f_k\|_2^2] \\ &= \frac{M_{2,k}}{N\rho_k}, \end{aligned} \quad (46)$$

which proves Equation (13).

D.2 Proof of Lemma 2

If $M_{2,k} = 0$, Equation (13) makes the claim immediate. Otherwise, apply Markov's inequality to the nonnegative random variable $\|\hat{g}_k - g_k\|_2^2$:

$$\begin{aligned} &\Pr \left(\|\hat{g}_k - g_k\|_2 > \sqrt{\frac{M_{2,k}}{\delta N \rho_k}} \right) \\ &= \Pr \left(\|\hat{g}_k - g_k\|_2^2 > \frac{M_{2,k}}{\delta N \rho_k} \right) \\ &\leq \frac{\mathbb{E}\|\hat{g}_k - g_k\|_2^2}{M_{2,k}/(\delta N \rho_k)} \\ &\leq \delta, \end{aligned} \quad (47)$$

where the last line uses Equation (13). Taking complements proves Equation (14).

D.3 Proof of Lemma 1

Let $\|\hat{g} - g\|_2 \leq \varepsilon$. Smoothness gives, for every $\gamma \geq 0$,

$$J(\theta + \gamma \hat{g}) - J(\theta) \geq \gamma g^\top \hat{g} - \frac{L\gamma^2}{2} \|\hat{g}\|_2^2. \quad (48)$$

The inner product satisfies

$$\begin{aligned} g^\top \hat{g} &= \|\hat{g}\|_2^2 + (g - \hat{g})^\top \hat{g} \\ &\geq \|\hat{g}\|_2^2 - \|g - \hat{g}\|_2 \|\hat{g}\|_2 \\ &\geq \|\hat{g}\|_2 (\|\hat{g}\|_2 - \varepsilon). \end{aligned} \quad (49)$$

Substituting Equation (49) into Equation (48) gives

$$J(\theta + \gamma \hat{g}) - J(\theta) \geq \gamma \|\hat{g}\|_2 (\|\hat{g}\|_2 - \varepsilon) - \frac{L\gamma^2}{2} \|\hat{g}\|_2^2. \quad (50)$$

For $\hat{g} \neq 0$, the derivative of the right-hand side is

$$\|\hat{g}\|_2 (\|\hat{g}\|_2 - \varepsilon) - L\gamma \|\hat{g}\|_2^2. \quad (51)$$

Its nonnegative maximizer is $\gamma = L^{-1}(1 - \varepsilon/\|\hat{g}\|_2)_+$. When $\hat{g} = 0$, both the chosen step and the claimed lower bound are zero. Substituting the maximizer into Equation (50) proves Equation (10).

D.4 Proof of Theorem 3

Lemma 2 shows that the event

$$\|\hat{g}_k - g_k\|_2 \leq \sqrt{\frac{M_{2,k}}{\delta N \rho_k}} \quad (52)$$

has probability at least $1 - \delta$. On this event, apply Lemma 1 with $\theta = \theta_k$, $\hat{g} = \hat{g}_k$, $\varepsilon = \sqrt{M_{2,k}/(\delta N \rho_k)}$, and $L = L_k$. The resulting inequality is exactly Equation (16).

D.5 Risk and certificate comparisons for coefficient rules

For a coefficient rule u , add and subtract the estimator mean:

$$\hat{g}_k(u) - g_k = \{\hat{g}_k(u) - \mathbb{E}[\hat{g}_k(u)]\} + \{\mathbb{E}[\hat{g}_k(u)] - g_k\}. \quad (53)$$

Squaring and taking expectation gives

$$\begin{aligned} m_k(u) &= \mathbb{E} \|\hat{g}_k(u) - \mathbb{E}[\hat{g}_k(u)]\|_2^2 + \|\mathbb{E}[\hat{g}_k(u)] - g_k\|_2^2 \\ &\quad + 2\{\mathbb{E}[\hat{g}_k(u)] - g_k\}^\top \mathbb{E}[\hat{g}_k(u) - \mathbb{E}[\hat{g}_k(u)]] \\ &= \mathbb{E} \|\hat{g}_k(u) - \mathbb{E}[\hat{g}_k(u)]\|_2^2 + \|\mathbb{E}[\hat{g}_k(u)] - g_k\|_2^2, \end{aligned} \quad (54)$$

because the centered term has mean zero. This proves Equation (17). For an iid sample mean, the first term in Equation (54) is the per-sample variation divided by N .

It remains to connect total risk to the expected policy-improvement certificate. For any random estimate $\hat{g}$ of g_k , smoothness gives

$$J(\theta_k + \eta_k \hat{g}) - J(\theta_k) \geq \eta_k g_k^\top \hat{g} - \frac{L_k \eta_k^2}{2} \|\hat{g}\|_2^2. \quad (55)$$

Taking expectation in Equation (55), with $\mathbb{E}\|\hat{g}\|_2^2 < \infty$, gives

$$\mathbb{E}[J(\theta_k + \eta_k \hat{g}) - J(\theta_k)] \geq \eta_k g_k^\top \mathbb{E}[\hat{g}] - \frac{L_k \eta_k^2}{2} \mathbb{E}\|\hat{g}\|_2^2. \quad (56)$$

Setting $\hat{g} = \hat{g}_k(u)$ proves Theorem 4. The polarization identity, followed by expectation, gives

$$2g_k^\top \mathbb{E}[\hat{g}] = \|g_k\|_2^2 + \mathbb{E}\|\hat{g}\|_2^2 - \mathbb{E}\|\hat{g} - g_k\|_2^2. \quad (57)$$

Substituting Equation (57) into the certificate and subtracting the IS certificate gives

$$\begin{aligned} B_k(u; \eta_k) - B_k(u_{\text{IS}}; \eta_k) &= \frac{\eta_k}{2} [m_k(u_{\text{IS}}) - m_k(u)] \\ &\quad - \frac{\eta_k}{2} (1 - L_k \eta_k) \{\mathbb{E}\|\hat{g}_k(u_{\text{IS}})\|_2^2 - \mathbb{E}\|\hat{g}_k(u)\|_2^2\}. \end{aligned} \quad (58)$$

Since $\eta_k > 0$, the right-hand side is positive exactly under Equation (20), proving Corollary 5. Since the IS estimator is unbiased, subtracting the certificate definitions instead gives

$$\begin{aligned} B_k(u; \eta_k) - B_k(u_{\text{IS}}; \eta_k) &= -\eta_k \left\{ \|g_k\|_2^2 - g_k^\top \mathbb{E}[\hat{g}_k(u)] \right\} \\ &\quad + \frac{L_k \eta_k^2}{2} \left\{ \mathbb{E}\|\hat{g}_k(u_{\text{IS}})\|_2^2 - \mathbb{E}\|\hat{g}_k(u)\|_2^2 \right\}. \end{aligned} \quad (59)$$

Its positivity is equivalent to Equation (21).

D.6 Intact prompt groups

Suppose M prompts are independent and each prompt has G conditionally independent responses. Keep each prompt group intact and define its contribution by $\Xi_{k,i} := G^{-1} \sum_{j=1}^G W_{k,ij} f_{k,ij}$. A leave-one-out baseline is admissible because it excludes the focal response. Conditional on the prompt and the remaining responses, the focal score still has mean zero. Hence $\mathbb{E}[\Xi_{k,i}] = g_k$.

Jensen’s inequality gives

$$\begin{aligned} \|\Xi_{k,i}\|_2^2 &= \left\| \frac{1}{G} \sum_{j=1}^G W_{k,ij} f_{k,ij} \right\|_2^2 \\ &\leq \frac{1}{G} \sum_{j=1}^G W_{k,ij}^2 \|f_{k,ij}\|_2^2. \end{aligned} \quad (60)$$

Repeating the sample-mean calculation across the M independent prompt groups gives

$$\begin{aligned} \mathbb{E} \left\| \frac{1}{M} \sum_{i=1}^M \Xi_{k,i} - g_k \right\|_2^2 &\leq \frac{1}{M} \mathbb{E} \|\Xi_{k,1}\|_2^2 \\ &\leq \frac{\|f_k\|_\infty^2}{M \rho_k}. \end{aligned} \quad (61)$$

Markov’s inequality then gives the bounded-contribution specialization of Lemma 2 with the number of independent prompts M in place of the number of responses N .

D.7 Bounded-contribution specialization

If $\|f_k\|_\infty < \infty$, then $M_{2,k} \leq \|f_k\|_\infty^2$. Lemma 2 therefore recovers the conventional radius $\|f_k\|_\infty / \sqrt{\delta N \rho_k}$ as a direct special case.

References

- [1] Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, Nicholas Joseph, Saurav Kadavath, Jackson Kernion, Tom Conerly, et al. Training a helpful and harmless assistant with reinforcement learning from human feedback, 2022.
- [2] Rasool Fakoor, Murdock Aubry, Nicholas Stranges, and Alexander J. Smola. Trust the batch, on- or off-policy: Adaptive policy optimization for RL post-training, 2026.
- [3] Luke J. Huang, Zhuoyang Zhang, Qinghao Hu, Shang Yang, and Song Han. Stable asynchrony: Variance-controlled off-policy RL for LLMs, 2026.
- [4] Chris Yuhao Liu, Liang Zeng, Jiacai Liu, Rui Yan, Jujie He, Chaojie Wang, Shuicheng Yan, Yang Liu, and Yahui Zhou. Skywork-Reward: Bag of tricks for reward modeling in LLMs, 2024.
- [5] Alberto Maria Metelli, Matteo Papini, Francesco Faccio, and Marcello Restelli. Policy optimization via importance sampling. In *Advances in Neural Information Processing Systems*, volume 31, 2018. URL <https://arxiv.org/abs/1809.06098>.

- [6] Alberto Maria Metelli, Matteo Papini, Nico Montali, and Marcello Restelli. Importance sampling techniques for policy optimization. *Journal of Machine Learning Research*, 21(141): 1–75, 2020. URL <https://jmlr.org/papers/v21/20-124.html>.
- [7] MiniMax. MiniMax-M1: Scaling test-time compute efficiently with lightning attention, 2025.
- [8] Matteo Papini, Matteo Pirotta, and Marcello Restelli. Smoothing policies and safe policy gradients. *Machine Learning*, 111:4081–4137, 2022. doi: 10.1007/s10994-022-06232-6. URL <https://doi.org/10.1007/s10994-022-06232-6>.
- [9] Jaesung R. Park, Junsu Kim, Gyeongman Kim, Jinyoung Jo, Sean Choi, Jaewoong Cho, and Ernest K. Ryu. Clip-low increases entropy and clip-high decreases entropy in reinforcement learning of large language models, 2025.
- [10] Matteo Pirotta, Marcello Restelli, and Luca Bascetta. Adaptive step-size for policy gradient methods. In *Advances in Neural Information Processing Systems*, volume 26, 2013. URL https://proceedings.neurips.cc/paper_files/paper/2013/hash/f64eac11f2cd8f0efa196f8ad173178e-Abstract.html.
- [11] Penghui Qi, Xiangxin Zhou, Zichen Liu, Tianyu Pang, Chao Du, Min Lin, and Wee Sun Lee. Rethinking the trust region in LLM reinforcement learning, 2026.
- [12] Alfréd Rényi. On measures of entropy and information. In *Proceedings of the Fourth Berkeley Symposium on Mathematical Statistics and Probability*, volume 1, pages 547–561. University of California Press, 1961. URL <https://projecteuclid.org/euclid.bsmsp/1200512181>.
- [13] John Schulman, Sergey Levine, Philipp Moritz, Michael I. Jordan, and Pieter Abbeel. Trust region policy optimization. In *Proceedings of the 32nd International Conference on Machine Learning*, volume 37 of *Proceedings of Machine Learning Research*, pages 1889–1897. PMLR, 2015. URL <https://proceedings.mlr.press/v37/schulman15.html>.
- [14] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms, 2017.
- [15] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. DeepSeekMath: Pushing the limits of mathematical reasoning in open language models, 2024.
- [16] Guobin Shen, Chenxiao Zhao, Xiang Cheng, Lei Huang, and Xing Yu. VESPO: Variational sequence-level soft policy optimization for stable off-policy LLM training, 2026.
- [17] Tim van Erven and Peter Harremoës. Rényi divergence and Kullback–Leibler divergence. *IEEE Transactions on Information Theory*, 60(7):3797–3820, 2014. doi: 10.1109/TIT.2014.2320500.
- [18] Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Weinan Dai, Tiantian Fan, Gaohong Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming Sheng, Yuxuan Tong, Chi Zhang, Mofan Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu, Jiaze Chen, Jiangjie Chen, Chengyi Wang, Hongli Yu, Yuxuan Song, Xiangpeng Wei, Hao Zhou, Jingjing Liu, Wei-Ying Ma, Ya-Qin Zhang, Lin Yan, Mu Qiao, Yonghui Wu, and Mingxuan Wang. DAPO: An open-source LLM reinforcement learning system at scale, 2025.
- [19] Haizhong Zheng, Jiawei Zhao, and Beidi Chen. Prosperity before collapse: How far can off-policy RL reach with stale data on LLMs?, 2025.